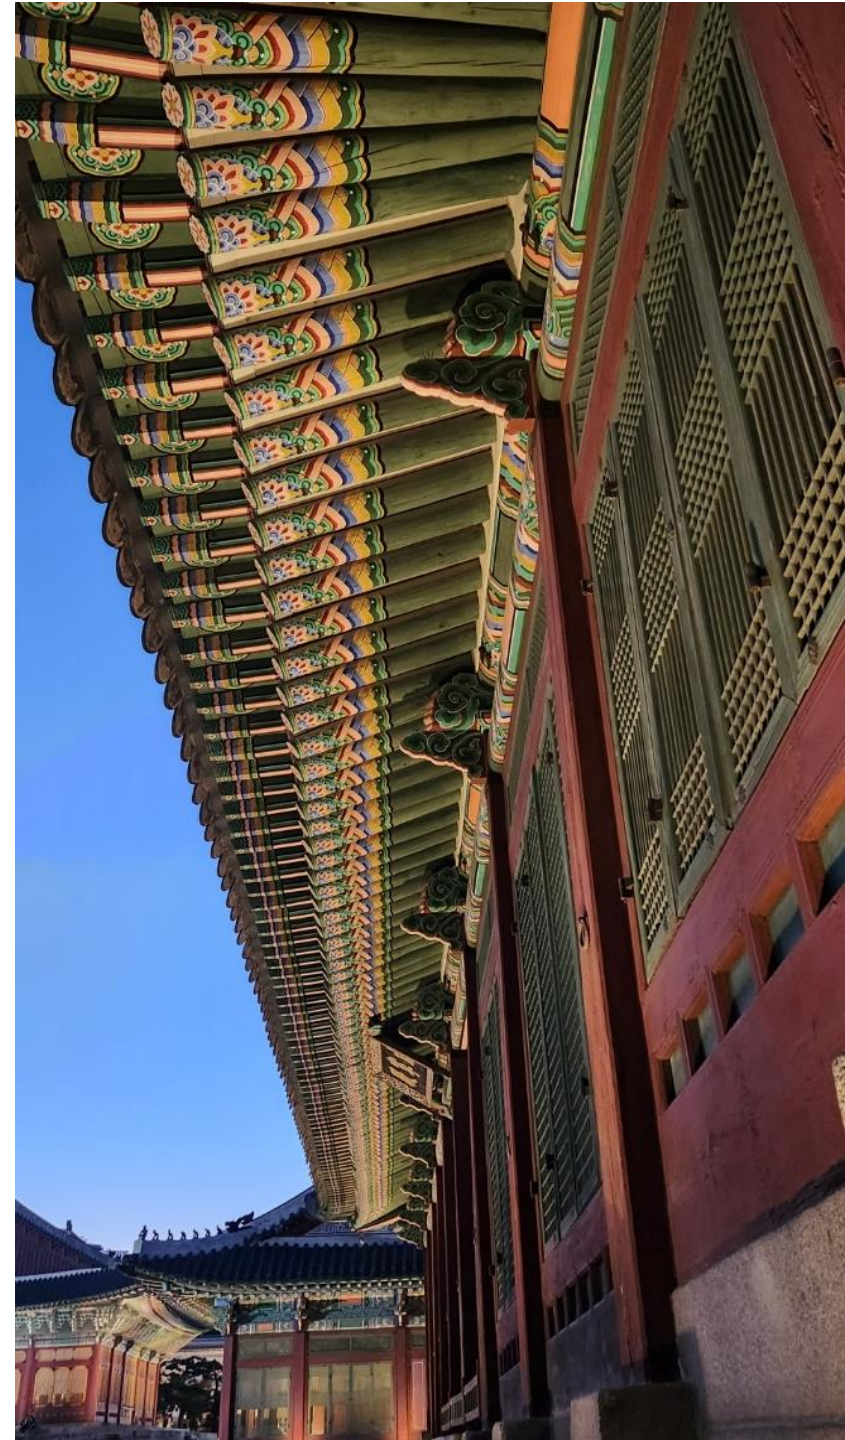


Linux File I/O

HGU



System call: a Cornerstone of System programming

- System call:
 - System programming starts and ends with *system calls*
 - **System calls** are function invocations made from user space – your text editors, favorite game, and so on to requests a service from the **kernel** (core part of OS)
- C library
 - On modern Linux systems, the C library is provided by **GNU** libc, *glibs*
 - The GNU C Library project provides the core library of GNU System
- C Compiler
 - In Linux, the standard C Compiler is provided by the **GNU Compiler Collection** (**gcc**)
- API (Application Programming Interface)
 - It is a software intermediary that allows two applications to talk to each other
 - It is based on source code

File Management Functions

- file descriptor
 - fd stands for **file descriptor**
 - File descriptor is an integer that uniquely identifies an opened file
 - A file descriptor is a non-negative integer, generally represented in the C programming language as the **type int**
- Most of the File Management Functions are declared in the `<fcntl.h>`, `<unistd.h>`, `<stdio.h>`, `<sys/stat.h>`, `<sys/types.h>`, and `<errno.h>`
 - `fcntl.h` (file control operation): `open`, `create`, `fcntl`
 - `unistd.h` (lower-level file handling): `read`, `write`, `close`, `lseek`, `unlink`
 - `stdio.h` (with `FILE` type): `fopen`, `fclose`, `fread`, `fwrite`, `fgets`, `fputs`, `frprintf`, `fscanf`
 - `sys/stat.h` (file information through `stat`): `stat`, `fstat`, `lstat`, `chmod`, `mkdir`
 - `sys/types.h` (data types used in system calls including file operations)

Opening Files

- A file is opened and a **file descriptor** is obtained with open() system call

- Syntax

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
#include <fcntl.h>
```

```
int open(const char * name, int flags); // open an existing file
```

```
int open (const char * name, int flags, mode_t mode);
```

```
    // create a new file if it does not exist when flags include O_CREAT
```

Opening Files using open()

- `int open (const char * name, int flags);`


File name with path specifies mode

- `int open (const char *name, int flags, mode_t mode);`


File name with path **O_CREAT** specifies mode in
3 octal digits for permission

open() flags

- **O_RDONLY** – open for reading only
- **O_WRONLY** – open for writing only
- **O_RDWR** -- open for reading and writing
- **O_CREAT** – create the file if it does not exist
- **O_EXCL** – when used with **O_CREAT**, if the file already exists it is an error and the `open()` will fail
- **O_APPEND** – open the file in append mode
- **O_SYNC**: Open the file for synchronous I/O. Writes will not return until the data is physically written to disk (all data and metadata are synchronized).
- **O_DSYNC**: Similar to **O_SYNC**, but only the file's data (not its metadata) is synchronized.
- **O_RSYNC**: Synchronizes reads with writes, ensuring that reads see the latest data written to the file.

Security of Files

- Linux is the Multi-user operating system which can be accessed by many users simultaneously
- This raises security concerns as an unsolicited user can corrupt, change or remove critical data
- For effective security, Linux divides authorization into two Levels:
 - Ownership
 - Permission

Ownership of Files

- Every file and directory on your Linux system is assigned 3 types of owner, given below:
- User: a user is the **owner of the file**. By default the person who created a file becomes its owner. (Hence, a user is also called an owner)
- Group
 - A user-group can contain multiple users. All users belonging to a group will have the same access permissions to the file
 - Instead of manually assigning permissions to each user, you could add all users to a group, and assign group permission to file
- Other
 - Any other user who has access to a file
 - Practically, it means everybody else.

Permissions of Files

- Every file and directory in your UNIX/Linux system has following 3 permissions defined for all the 3 owners discussed above
- The three permissions are
 - Read: permission giving the authority to **open and read** a file
 - Write: permission giving the authority to **modify** the contents of a file
 - Execute: permission giving the authority to **execute** a file

sample:

File Permission by ls -l command shows 9 characters

-rwxrwxrwx

or 3 octal numbers
for the mode arg of open()

file type
d, l, c, b

user group others

```
# ls -l file
-rw-r--r-- 1 root root 0 Nov 19 23:49 file
```

File type

Owner (rw-)
Group (r--)
Other (r--)

r = Readable
w = Writeable
x = Executable
- = Denied

Permission of Files (Open System call mode)

- S_IRWXU : (00700) – All authority for owner
- S_IRUSR : (00400) – Read authority for owner
- S_IWUSR : (00200) – Write authority for owner
- S_IXUSR : (00100) – Execute authority for owner
- S_IRWXG : (00070) – All authority for group
- S_IRGRP : (00040) – Read authority for Group
- S_IWGRP : (00020) – Write authority for Group
- S_IXGRP : (00010) – Execute authority for Group
- S_IRWXO : (00007) – All authority for Others
- S_IROTH : (00004) – Read authority for Others
- S_IWOTH : (00002) – write authority for Others
- S_IXOTH : (00001) – Execute authority for Others

Return Values

- Both `open()` and `create()`
 - On Success, returning zero(**0**)
 - On Error, returning a non-zero number (**-1**), and **set errno** to an appropriate error value

open() Sample Code

```
int fd;  
fd = open ("home/user1/car", O_RDONLY);  
if (fd == -1) /* error */
```

the file /home/usr1/car .is opened for reading

```
int fd;  
fd = open ("home/user1/pearl," O_WRONLY | O_TRUNC);  
if (fd == -1) /* error */
```

if the file already exists, it will be truncated to a length of zero.
If the file does not exists, the call will fail.

```
int fd;  
fd = open (file, O_WRONLY | O_CREAT | O_TRUNC ,  
           S_IWUSR | S_IRUSR | S_IWGRP | S_IRGRP | S_IROTH);  
if (fd == -1)  
    /* error */
```

If the file does not exist, assuming a umask of 022 it is created with the permission 0644.
If it exists, it is truncated to zero length

rw-rw-r-- → 0661

0661 & ~0022 = 0644

create()


- create() System call create and open a file
- create is a special case of open() system call
- Syntax:

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
#include <fcntl.h>
```

```
int creat (const char *name, mode_t mode);
```


File name with path specifies mode

create() sample code

```
int fd;
```

```
fd = create( file, 0644);  
if(fd == -1 )  
    /* error handling */
```

==

```
int fd;
```

```
fd = open( file, O_WRONLY | O_CREAT | O_TRUNC, 0644);  
if(fd == -1 )  
    /* error handling */
```

Reading via read()

- The mechanism used for reading : the read() system call
- Syntax:

#include <unistd.h>

ssize_t read(int fd, void * buf, size_t len);

File descriptor

address of
the first byte
of buffer

length or count of buffer

- The **ssize_t** type is *signed version of size_t* (the negative values are used to connect errors)
- the maximum value of an ssize_t is **SSIZE_MAX** which is 2,147,483,647 bytes on a 32-bit machine

Reading via read()

- When read() returns successfully, its return value is **the number of bytes** actually read and placed in the buffer
- if len is zero, read returns zero and has no other results
- On Success, a non-negative integer is returned indicating the number of bytes actually read
- Otherwise(Failure), a **-1** is returned

Call to read() Results in many possibilities

- Case 1: returns a value equal to len. The results are as intended(**Success**).
- Case 2: returns a value less than len, but greater than 0 if any *interrupt* occurs during read or EOF is reached before len bytes were read.
- Case 3: **returns 0**, this indicates reaching *EOF and* there is nothing to read
- Case 4: the call blocks because no data is currently available. (It won't happen in nonblocking mode)
- Case 5: returns -1
 - `errno == EINTR` : a signal was received before any bytes were read. The call can be reissued.
 - `errno == EAGAIN` : read would block because no data is currently available and the request should be reissued later. It happens *only in nonblocking mode*
 - Other `errno` error values : serious error (EINVAL, EBADF, EFAULT, EIO, ENOMEM. ...)

read() Sample Code

```
unsigned long word;
ssize_t nr;

/* read a couple bytes into 'word' from 'fd' */
nr = read (fd, &word, sizeof(unsigned long));
if( nr == -1)
    /* error handling*/
```

```
/* reading all bytes */
ssize_t ret;

while ( len != 0 &&
        (ret = read (fd, buf, len)) != 0) {
    if(ret == -1){
        if(errno == EINTR)
            continue;
        perror("read");
        break;
    }
    len -= ret;
    buf += ret;
}
```

Nonblocking Reads

- By default, *read()* waits until at least one byte is available to return to the application; this default is called “**blocking**” mode
- Alternatively, individual file descriptors can be switched to “**non-blocking**” mode, which means that a *read()* on a slow file will return immediately, even if no bytes are available.
- Nonblocking I/O, if given file descriptor was opened in nonblocking mode.

← **open()** with **O_NONBLOCK** flag

Other Error Values of read()

- Possible errno values after a failure on read() include:
 - EBADF: The given file descriptor is invalid or is not open for reading
 - EFAULT: The pointer provided by buf is not inside the calling process's address space
 - EINVAL: The file descriptor is mapped to an object that does not allow reading
 - EIO: a low-level I/O error occurred
 - ENOMEM: not enough memory is available

Writing with write()

- The most basic and common system call for writing : write()
- it writes data from a buffer declared by the user to a given device, such as file
- This is the primary way to output data from a program by directly using a system call
- Syntax:

```
#include <unistd.h>
```

```
ssize_t write (int fd,    const void * buf,    size_t count);
```

File
descriptor

address
of source

length of the data to be
written

Return value of write()

- On success, the number of bytes written
- On Error, -1 is returned and errno is set appropriately
- Partial writes:
 - a successful write() may transfer fewer than count bytes
 - Such partial writes can occur for various reasons
 - If there was insufficient space to the disk to write all of the requested bytes
 - Because of interrupt

write() mode

- Append Mode
 - When fd is *opened in append mode* (**O_APPEND**), writes do not occur at the fd's current file position. Instead, they occur at the current end of the file
- Nonblocking Writes
 - When fd is *opened in nonblocking* mode (via **O_NONBLOCK**), and the write as issued would normally block, the write() system call returns -1 and set errno to EAGAIN
 - The request should be reissued later

Other Error Codes of Write

- EBADF: The given file descriptor is not valid or is not open for writing
- EFAULT: the pointer provided by buf points outside of the process's address space
- EINVAL: The given file descriptor is mapped to an object that is not suitable for writing
- ENOSPC: the file system backing the given file descriptor does not have sufficient space
- EINTR: write was blocked by interrupt (in this case retry write())

simple Write() sample

```
ssize_t ret, nr;

while (len != 0 && (ret = write(fd, buf, len))!= 0) {
    if(ret == -1) {
        if(errno == EINTR)
            continue;
        perror("write");
        break;
    }
    len-= ret;
    buf += ret;
}
```


How to perform write()

- Size Limits on write()
 - If count is larger than SSIZE_MAX, the results of the call to write are undefined
 - A call to write() with count of zero results in the call returning immediately with a return value of 0
- Behavior of write
 - When a user-space application issues a write() system call, the Linux kernel performs a few checks and then simply *copies the data into a buffer*
 - Later, in the background, the kernel gathers up all of the dirty buffers, *which are buffer that contain data newer* than what is on disk, sorts them optimally, and writes them out to disk.

Synchronized I/O

user issues write() some data which contains three lines to file called sample.txt located in your z drive

request



Kernel or OS accept the request and collect data from write

time frame 1: data is collected to buffer
time frame 2: data is collected to buffer
time frame 3: data is collected to buffer
and then

It will sort the data in buffers and in optimal time the consolidated data will be moved to sample.txt located in your z drive

process

Synchronized I/O

- By default, the Linux kernel writes data to disk asynchronously
- Writes are buffered (cached) in memory, and written to the storage device at the optimal time
- The **Synchronous I/O** provides some functions to ensure that all operations finish before they return.
- **Sync** System call:
 - The sync system call forces an immediate write of all cached data to disk but it doesn't wait to complete
 - This call initiates the process of committing all buffers to disk
- Sync Syntax: `sync [option] [file]`
- Examples:
 - `sync -d` : sync only file data not metadata
 - `sync -f <file>`: sync file systems which contains the files

Synchronized I/O: fsync()

- Syntax of fsync()

#include <unistd.h>

int fsync(int fd);

fd: file descriptor of open file

- The call to above function ensures that *all dirty data* associated with the file mapped by the file descriptor fd are written back to disk
- fsync() writes back both **data and metadata (information summary)**

Synchronized I/O: fdatasync()

- Syntax

```
#include <unistd.h>
```

```
int fdatasync(int fd);
```

- A call to fdatasync() ensures that only the file's data blocks are synchronized to disk, along with essential metadata like file size.
- The call does **not** guarantee that nonessential metadata such as access time, ownership, or permission is synchronized to disk,
- is therefore potentially faster

Flags of open () for Synchronized I/O

- **O_SYNC** flag requires that any write operations block until all data and all metadata have been written to persistent storage
- **O_DSYNC** flag specifies that only normal data be synchronized after each write operation, not metadata

Direct I/O

- The Linux kernel implements a complex layer of caching, buffering, and I/O management between devices and applications
- A high-performance application may wish to bypass this layer of complexity and perform its own I/O management
- Providing the **O_DIRECT** flag to **open()** instructs the kernel to minimize the presence of I/O management
 - When O_DIRECT flag is provided, I/O will initiate directly from user-space buffer to the device, *bypassing the page cache*
 - Direct I/O will be synchronous; operations will not return until completion
- For an I/O operation to be performed as direct I/O, it must meet certain alignment criteria (depending on disk driver, disk controller, and memory management hardware). Otherwise the direct I/O request is performed as data synchronous I/O.

Closing Files

- after a program has finished working with a file descriptor, it can unmap the file descriptor from the associated file via `close()` system call:

```
#include <unistd.h>
```

```
int close(int fd);
```

- A call to `close()` unmaps the open file descriptor `fd` and disassociates the file from process
- It is a common mistake to not check the return value if `close()`
- There are handful of possible `errno` values on failure. Other than `EBADF` the most important error value is `EIO`, indicating a low-level I/O error probably unrelated to the actual close

Seeking : lseek()

- lseek() is a system call that is use to *change the location of the read/write pointer* of a file descriptor
- Syntax:

```
#include <sys/types.h>
#include <unistd.h>
off_t lseek(it fd, off_t pos, int origin);
```
- origin parameter
 - **SEEK_CUR** : the current file position of fd is set to its current value plus pos, which can be negative, zero, or positive. **A pos of zero returns current file position**
 - **SEEK_END**: the current file position of fd is set to the current length of file + pos. **A pos of zero sets the offset to the end of the file**
 - **SEEK_SET**: the current file position of fd is set to pos. **A pos of zero sets the offset to the beginning of the file.**

Seeking: lseek() Error Values

- EBADF: the given *file descriptor* does not refer to an open file descriptor
- EINVAL: the value given for *origin* is not valid or the resulting file position would be negative
- EOVERFLOW: *The resulting file offset cannot be represented in an off_t*
- ESPIPE: the given file descriptor is associated with an *unseekable object*, such as a pipe, FIFO, or socket

Positional Reads/Writes

- Instead of lseek(), Linux provides two variants of the read() and write() system calls
- Both receive the file position from which to read or write
- Upon completion, they do *not update the file position*
- The read form is called pread():

```
#define _XOPEN_SOURCE 500
#include <unistd.h>
ssize_t pread(int fd, void *buf, size_t count, off_t pos);
```
- This call read up to *count* bytes into buf from the file descriptor *fd* at file position *pos*.

Positional Reads/Writes

- The write form is called *pwrite()*:
#define _XOPEN_SOURCE 500
#include <unistd.h>
ssize_t **pwrite**(int fd, void *buf, size_t count, **off_t pos**);
- This call writes up to count bytes from buf to the file descriptor fd at file position pos.

Truncating Files

- Linux provides two system calls, `ftruncate()` and `truncate()`, for truncating the length of a file
- Syntax of `ftruncate`:
`#include <unistd.h>`
`#include <sys/types.h>`
`int ftruncate(int fd, off_t len);`
- Syntax of `truncate`:
`#include <unistd.h>`
`#include <sys/types.h>`
`int truncate(const char *name, off_t len);`

Truncating Files

- These system calls usually truncate a file to a *size smaller than its current length*
- the **ftruncate()** system call operates on the file descriptor given by ***fd***, which must be open for writing
- The **truncate()** system call operates on the ***filename*** given by path, which must be writable.
- Returns
 - On success, return **0**
 - On error, they return **-1** and set **errno** as appropriate

Kernel Internals

- The Kernel subsystem consists of :
 - the virtual filesystems (VFS)
 - the page cache, and
 - page writeback
- *The virtual filesystem*, occasionally called a virtual file switch, is a mechanism of *abstraction* that allows the Linux kernel to call **filesystem** functions and manipulate filesystem data without knowing the specific type of filesystem being used
- *Linux filesystem* is generally a *built-in layer of a Linux OS* used to handle the data management of the storage.
 - helps to arrange the file on the disk storage.
 - manages the file name, file size, creation date, and more information about a file

Kernel Internals


- The Page Cache:
 - the page cache is an in-memory store of recently accessed data from an on-disk filesystem.
 - Disk access is painfully slow, particularly relative to today's processor speeds
 - Storing requested data in memory allows the kernel to fulfill subsequent requests for the same data from memory, avoiding repeated disk access

Kernel Internals

- Page Writeback
 - Eventually the dirty buffers need to be committed to disk, synchronizing the on-disk files with the data in memory. This is known as writeback. It occurs in two situations:
 - When free memory shrinks below a configurable threshold, dirty buffers are written back to disk so that the now-clean buffers may be removed, freeing memory
 - When a dirty buffer ages beyond a configurable threshold, the buffer is written back to disk. This prevents data from remaining dirty indefinitely

Buffered I/O

HGU



Introduction on buffered I/O

- I/O **buffering** is a mechanism that improves the throughput of I/O operation.
 - Throughput is the amount of work completed in a unit of time
- It is implemented directly in hardware and the corresponding drivers (hence the block devices found in UNIX-like systems), and also universal among programming language standard libraries

I/O Latencies and Buffering

- I/O operations often have high latencies.
 - latency: the time between the initiation of I/O process and its completion
 - Most of this latency is due to hardware itself; for example, information cannot be read from or written to a hard disk until spinning of the disk brings target sectors under the disk head
- The Latency is reduced by having one or more input and output buffers associated with each device
- A *buffer* is a memory area that stores data being transferred between two devices or between a device and an application

I/O Buffering

- The goal of the buffering provided by the standard I/O library
- Buffering is done for three reasons
 - To cope with a **speed mismatch** between producer and consumer of a data stream
 - To provide adaptation for data that have **different data-transfer sizes**
 - To support **copy semantics** for the application I/O
- Copy semantics
 - Data is first copied from user application memory into kernel memory
 - Data from kernel memory is then written to device
 - Prevents application from changing contents of a buffer before it is done being written

User-Buffered I/O

- User buffered I/O, shortened to *buffering or buffered I/O*, refers to the technique of temporarily *storing the results of an I/O operation in user-space* before transmitting it to the kernel (in case of write) or before providing it to user process (in the case of reads)
- By so buffering the data, you can minimize the number of system calls and can block-align I/O operations, which may *improve the performance* of user application.

Block Size and Buffered I/O

- In practice, Block size are 512, 1024, 2,048, 4,096, or 8,192 bytes.
- A large performance gain is realized simply by performing operations with respect to block size
- This is because the kernel and hardware speaks in terms of blocks
- Thus, using block size or a value that fits neatly inside of a block guarantees block-aligned I/O and prevents extraneous (external) work inside kernel.

Example(1): A Simple getchar()

```
int getchar(void) {  
    static char c;  
    if (read(0, &c, 1) == 1)  
        return c;  
    else return EOF;  
}
```

Read one character from `stdin`

- File descriptor 0 is `stdin`
- `&c` points to the buffer
- 1 is the number of bytes to read

Read returns the number of bytes read

- In this case, 1 byte means success


Example(2): Making getchar() More Efficient

- Poor performance reading one byte at a time
 - read() system call is accessing the device (e.g., a disk)
 - Reading one byte from disk is very time consuming
 - Better to read and write in *larger chunks*
- Buffered I/O
 - Read a large chunk from disk into a buffer
 - Similarly, for writing, write individual bytes to a buffer
 - And write to disk when full, or when stream is closed
 - Known as “flushing” the buffer

Example(3): Better getchar() with Buffered I/O

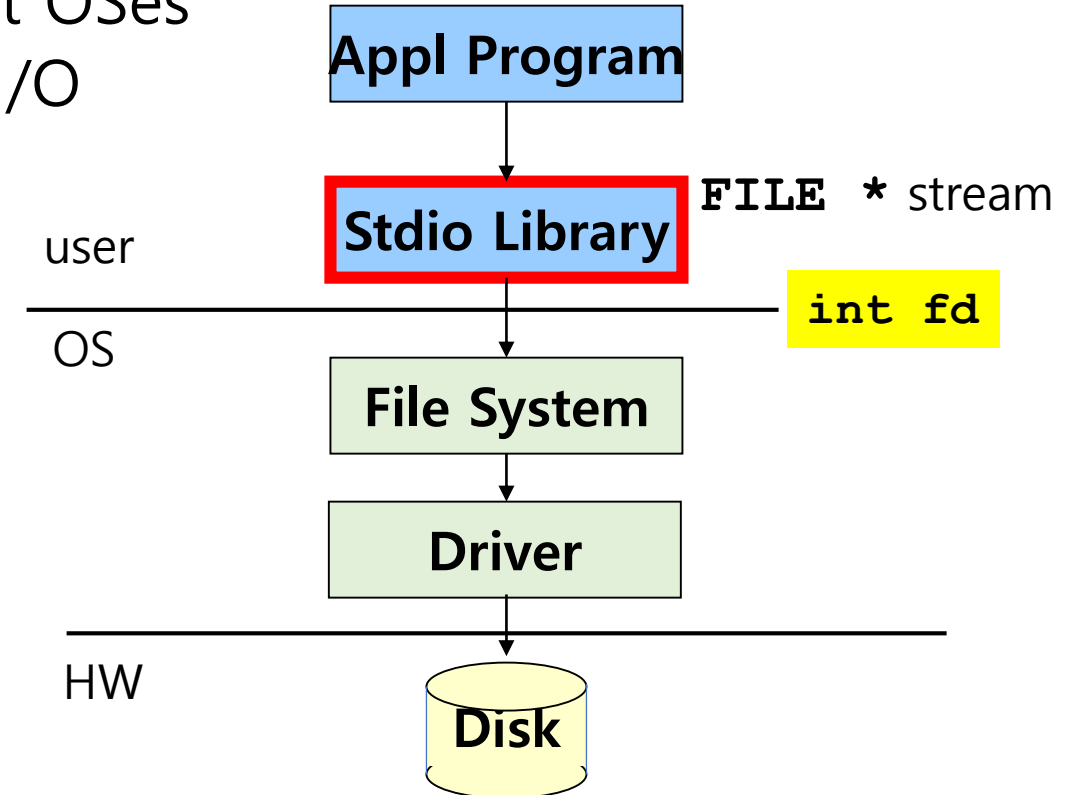
```
int getchar(void) {  
    static char base[1024];  
    static char *ptr;  
    static int cnt = 0;  
  
    if (cnt--)  
        return *ptr++;  
  
    cnt = read(0, base, sizeof(base));  
    if (cnt <= 0)  
        return EOF;  
    ptr = base;  
    return getchar();  
}
```

} persistent variables



Standard I/O Library

- Portability
 - Generic I/O support for C programs
 - Specific implementations for various host OSes
 - Invokes the OS-specific system calls for I/O
- Abstractions for C programs
 - Streams
 - Line-by-line input
 - Formatted output
- Additional optimizations
 - Buffered I/O
 - Safe writing



The Core of Standard I/O Library in C


1. FILE Structure: fd, buffer, permission
2. Standard Streams: stdin(0), stdout(1), stderr(2)
3. FILE I/O functions
 - Generic: fopen(fp), fclose(fp), fread(buf, size, count, fp), fwrite(buf, size, count, fp),
 - Text file I/O: fgets(buf, size, fp), fputs(buf, fp)
4. Formatted I/O: printf, scanf, fprintf(fp, format, ...), fscanf(fp, format, ...)
5. character based I/O: fgetc(fp), fputc(c, fp), getchar(), putchar(c)
6. Buffering: fflush(fp)
7. Error Handling : ferror(fp), feof(fp), clearerr(), perror()

Stream Abstraction

- Any source of input or destination for output is regarded as file
 - E.g., keyboard as input, and screen as output
 - E.g., files on disk or CD, network ports, printer port, ...
- Accessed in C programs through **file pointers**
 - E.g., `FILE *fp1;`
 - E.g., `fp1 = fopen("myfile.txt", "r");`
- The standard C library provides I/O library (**stdio**), which in turn provides a platform-independent user-buffering solution
 - **Three standard streams** provided by stdio.h
 - Standard Streams: **stdin**, **stdout**, and **stderr**

Standard I/O Streams

- The three I/O connections are called **stdin (standard input)**, **stdout (standard output)**, and **stderr (standard error)**.
- stdin is a stream from which a program reads its input data
- stdout is a stream to which a program writes its output data
- stderr is another output stream typically used by programs to output error messages or diagnostics.



Sequential Access to a Stream

- Each stream has an associated **file position**
 - Starting at beginning of file (if opened to read or write)
 - Or, starting at end of file (if opened to append)



- **Read/write** operations **advance the file position**
 - Allows sequencing through the file in sequential manner
- Support for random access to the stream
 - Functions to learn current position and **seek** to new one

Standard I/O: File Pointer

- Standard I/O Routines do not operate directly on *file descriptors*.
- Instead, they use **file pointer**
- file pointer is a pointer which is used to handle and keep track on the files being accessed
- a new data type called "FILE" is used to declare file pointer
FILE *fp; // fp is a file pointer
- **FILE** is defined in stdio.h

Details of FILE in stdio.h (K&R 8.5)

```
#define OPEN_MAX 20    /* max files open at once */

typedef struct _iobuf {
    int  cnt;          /* num chars left in buffer */
    char *ptr;         /* ptr to next char in buffer */
    char *base;        /* beginning of buffer */
    int  flag;         /* open mode flags, etc. */
    char fd;           /* file descriptor */
} FILE;

extern FILE _iob[OPEN_MAX];

#define stdin (&_iob[0])
#define stdout (&_iob[1])
#define stderr (&_iob[2])
```

Opening File with fopen()

- files are opened for reading or writing via fopen():

#include <stdio.h>

FILE *fopen(const char *file_path, const char *mode);

- This function opens the file path with the behavior given by mode and associates a new **stream** with it
- A stream is a sequence of data elements made available over time

Mode of fopen()

- **r** : open file for reading. stream is positioned at the start of the file
- **r+** : open file for both reading and writing. stream is positioned at the start of the file
- **w** : open the file for writing. If file exists, it is truncated to zero length. If the file does not exist it is created.
- **w+** : open the file for both reading and writing. If file exists, it is created.
- **a** : open the file for writing in append mode. The file is created if it does not exist. The stream is positioned at the end of the file.
- **a+** : open the file for both reading and writing in append mode. The file is created if it does not exist. The stream is positioned at the end of the file.

fopen vs. open

fopen()	open()
fopen series are standard C library functions	the open series are defined by POSIX and are system calls in UNIX system
when using fopen functions, you must define an object that refers to a file. It is called " file handler " and is a struct	the Open series uses an int integer call " file descriptor "
On failure, returns NULL and sets errno On success, returns valid pointer to FILE	On failure, returns -1 and sets errno On success, returns file descriptor (int)

Opening a Stream via File Descriptor

- function `fdopen()` converts an already open file descriptor to a stream
`FILE *fdopen(int fd, const char *mode);`
- The possible modes are the same as for `fopen()` must be compatible with the modes originally used to open the file descriptor

```
FILE * stream;
int fd;
fd = open("/home/kid/map.txt", O_RDONLY);
if (fd == -1)
    /* Error Handling */
stream = fdopen(fd, "r");
if(stream == NULL)
    /* Error Handling */
```

Closing Streams with fclose()

- **fclose()** function closes a given stream:
#include <stdio.h>
int fclose(FILE *fp);
- Any buffered and not-yet-written data is first flushed.
- On success, fclose() returns 0.
- On Failure, it returns EOF and sets errno appropriately.
- **fcloseall()** function closes all streams associated with current process including stdin, stdout, and stderr:
#define _GNU_SOURCE
#include <stdio.h>
int fcloseall(void);
/ fcloseall always returns 0 */*

Reading from Stream

- standard C library implements multiple functions for reading from an open stream
- most popular 3 functions:
 - reading one character at a time: **fgetc()**
 - reading one entire text line at a time: **fgets()**
 - reading binary data: **fread()**

Reading a Character: fgetc()

- The ideal I/O pattern is simply reading one character at a time. The function fgetc() is to read a single character from a stream

```
#include <stdio.h>
```

```
int fgetc(FILE *stream);
```

- This function reads the next character (**unsigned char**) from the specified stream and advances the position indicator for the stream

fgetc() Sample Code

```
#include <stdio.h>
int main() {
    FILE *fp;
    int c;
    int n = 0;
    fp = fopen("file.txt", "r");
    if(fp==NULL) {
        perror("Error in opening file");
        return -1;
    }
```

```
do {
    c = fgetc(fp);
    if( feof(fp))
        break;
    printf("%c", c);
} while(1);
fclose(fp);
return(0);
}
```

Putting the character back : ungetc()

- Standard I/O provides a function for pushing a character back onto a stream, allowing you to “peek” at the stream and return the character if it turns out that you don’t want it:

```
#include <stdio.h>
```

```
int ungetc(int c, FILE *stream);
```

ungetc() Sample Code

```
#include <stdio.h>
int main() {
    FILE *fp;
    int c;
    char buffer[256];

    fp = fopen("file.txt", "r");
    if(fp==NULL) {
        perror("Error in opening file");
        return (-1);
    }
```

```
while (!feof(fp)) {
    c = fgetc(fp);
    if(c == '!') /* replace ! with + */
        ungetc('+', fp);
    else
        ungetc(c, fp);
    c= fgetc( fp);
    fputc(c, stdout);
}
return(0);
}
```

Reading an Entire Line: fgets()

- the function fgets() reads a string from a given stream:

`#include <stdio.h>`

`char *fgets(char *str, int size, FILE *stream)`

- This function reads up to one less than size bytes from stream and stores the results in str
- a null character('\0') is stored in the buffer after the last byte read in

fgets() Sample Code

```
#include <stdio.h>
int main() {
    FILE *fp;
    char str[60];
    /* opening file for reading */
    fp = fopen("file.txt", "r");
    if (fp == NULL) {
        perror("Error opening file");
        return (-1);
    }
}
```

```
if ( fgets(str, 60, fp) != NULL){
    /* writing contents to stdout*/
    puts(str);
}
fclose(fp);
return (0);
}
```

Reading Binary Data: fread()

- Sometimes, developers want to read and write complex binary data
- For this, the standard I/O library provides **fread()**:
 - reads data from a file and store it in a buffer

#include <stdio.h>

size_t fread(void *buffer, size_t size, size_t count, FILE *stream);

- buffer: pointer to the buffer where data will be stored. A buffer is region of memory used to temporarily store data
- size: the size of each elements to be read
- count: the number of elements to be read
- stream: pointer to the FILE object from where data is to be read
- Upon Success execution it returns integer value equivalent to count. In case of an error or EOF, a value less than count is returned

Writing to a Stream

- standard C library implements multiple functions for writing to an open stream
- most popular 3 functions:
 - writing one character at a time: `fputc()`
 - writing one entire line at a time: `fputs()`
 - writing binary data: `fwrite()`

Writing a Character: fputc()

- The counterpart of function fgetc() is fputc()

#include <stdio.h>

int fputc(int c, FILE *stream);

- Parameters
 - c : This is the character to be written. This is passed as its int promotion
 - stream : it is the pointer to a FILE object that identifies the stream where the character is to be written
- Return Value
 - If there are no errors, the same character that has been written is returned. If Error occurs, EOF is returned and the error indicator is set.

fputc() Sample Code

```
#include <stdio.h>
int main() {
    FILE *fp;
    int ch;

    fp = fopen("file.txt", "w+");
    if(fp==NULL) {
        perror("Error in opening file");
        return -1;
    }
```

```
    for(ch = 33; ch <= 100; ch++)
        fputc( ch, fp);
    fclose(fp);
    return(0);
}
```

Writing a String of Characters: fputs()

- the function fputs() writes an entire string to a given stream:

```
#include <stdio.h>
```

```
int fputs(char *str, FILE *stream)
```

- Parameters
 - str: an array containing the null-terminated sequence of characters to be written
 - stream: the pointer to a FILE object that identifies the stream where the string is to be written
- Return Value: the function returns a non-negative value, or else on error it returns EOF

fputs() Sample Code

```
#include <stdio.h>
int main()
{
    const char *buffer = "Hello world!";
    fputs (buffer, stdout);
    return (0);
}
```

Writing Binary Data: fwrite()

- To directly store binary data such as C variables, standard I/O provides fwrite():

#include <stdio.h>

size_t fwrite(void *buffer, size_t size, size_t nr, FILE *stream);

- buffer: pointer to the array of elements to be written
 - size: the size in bytes of each elements to be written
 - count: the number of elements, each one with the size of **size** bytes
 - stream: pointer to the FILE object that specifies an output stream
- Return value: It returns the total number of elements successfully returned as a size_t object, which is an integral data type.

fwrite() Sample Code

```
#include <stdio.h>
int main()
{
    FILE *fp;
    int iCount;
    char arr[6] = "Hello";

    fp = fopen("sample.txt", "wb");
    iCount = fwrite(arr, 1, 5, fp);
    fclose(fp);
    return (0);
}
```

Complete Sample Code for Buffered I/O

```
#include <stdio.h>

int main (void) {
    FILE *in, *out;
    struct pirate {
        char name[100]; /* real name */
        unsigned long booty; /* in pounds sterling */
        unsigned int beard_len; /* in inches */
    } p, blackbeard = { "Edward Teach", 950, 48 };
    out = fopen ("data", "w");
    if (!out) {
        perror ("fopen");
        return 1;
    }
    if (!fwrite (&blackbeard, sizeof (struct pirate), 1, out)) {
        perror ("fwrite");
        return 1;
    }
}
```

```
    if (fclose (out)) {
        perror ("fclose");
        return 1;
    }
    strcpy(p.name,"None"); p.booty = 0L; p.beard_len = 0;
    in = fopen ("data", "r");
    if (!in) {
        perror ("fopen");
        return 1;
    }
    if (!fread (&p, sizeof (struct pirate), 1, in)){
        perror ("fread");
        return 1;
    }
    if (fclose (in)) {
        perror ("fclose");
        return 1;
    }
    printf ("name=%sW" booty=%lu beard_len=%uWn",
           p.name, p.booty, p.beard_len);
    return 0;
}
```

Seeking a Stream: fseek()

- fseek() function, the most common standard I/O seeking interfaces, manipulates the file position of stream in accordance with offset where:

#include <stdio.h>

int fseek(FILE *stream, long offset, int whence);

- Parameters
 - stream: the pointer to FILE object that identifies the stream
 - offset: the number of bytes to offset from whence
 - whence: the position from where offset is added


whence value	SEEK_SET	SEEK_CUR	SEEK_END
description	beginning of file	current position of file	end of file

- Return value: On success Zero(0), On failure non-zero value

fseek() Sample Code

- random search

```
#include <stdio.h>
int main() {
    FILE *fp = fopen("file.txt", "w+");
    fputs("This is tutorials-point.com", fp);
    fseek (fp, 7, SEEK_SET);
    fputs(" C Programming Language", fp);
    fclose(fp);
    return (0);
}
```



This isC Programmingint.com

Seeking a Stream

- Alternative to `fseek()` is **`fsetpos()`** and **`rewind()`**

- `fgetpos()`:

`#include <stdio.h>`

`int fsetpos(FILE *stream, fpos_t *pos);`

- Upon Success, it returns 0, and places the current stream position of stream in pos

- `rewind()`:

`#include <stdio.h>`

`void rewind (FILE *stream);`

- sets the position back to the start of the stream
- equivalent to `fseek(stream, 0, SEEK_SET)`

Obtaining the Current Stream Position: ftell()

- Unlike lseek(), fseek() does not return the updated position.
- the **ftell()** function returns the current stream position in the stream:

```
#include <stdio.h>
```

```
long ftell (FILE *stream);
```

- On Error, ftell() returns -1 and errno is set appropriately

```
#include <stdio.h>
int main() {
    int len;
    FILE *fp = fopen("file.txt", "w+");
    if(fp == NULL) perror("Error open");
    fseek (fp, 0, SEEK_END);
    len = ftell(fp);
    fclose(fp);
    printf("Size of file : %d bytes", len);
    return (0);
}
```

Flushing a Stream: fflush()

- The standard I/O library provides an interface for writing out the user buffer to the kernel, ensuring that all data written to a stream is flushed via write(). The fflush() functions provides this functionality:

```
#include <stdio.h>
```

```
int fflush(FILE *stream);
```

- On invocation, any unwritten data in the stream pointed to by stream is flushed to the kernel.
- If stream is NULL, **all open input stream in the process are flushed**
- **On success, fflush returns 0, On failure it returns EOF and errno is set**

Errors and EOF

- Some of standard I/O interfaces, such as `fread()`, communicate failure back to the caller poorly, as they provides no mechanism for differentiating between error and end-of-file(EOF)
- With these calls, and on other occasions, it can be useful to check the status of a given stream to determine *whether it has encountered an error or reached end-of-file*.
- *Standard I/O provides two interfaces to this end.*
 - *`ferror()`*
 - *`feof()`*

Errors and EOF

- The function `ferror()` tests whether the error indicator is set on stream:
`#include <stdio.h>`
`int ferror(FILE *stream);`
- The error indicator is set by standard I/O interfaces in response to an error condition
- The function **returns *a nonzero value* if the indicator is set**, and **0** otherwise

Errors and EOF

- The function `feof()` tests whether the EOF indicator is set on stream:
`#include <stdio.h>`
`int feof(FILE *stream)`
- The EOF indicator is set by standard I/O interfaces when the end of file is reached
- The ***function returns nonzero*** value if the indicator is set, and **0** otherwise (not EOF)

Errors and EOF

- The `clearer()` function clears the error and the EOF indicator for stream:

```
#include <stdio.h>
```

```
void clearer(FILE *stream);
```

- It has no return value, and cannot fail
- You should make a call to `clearer()` only after checking the error and EOF indicator

Obtaining the File Descriptor Associated

- To obtain the file descriptor backing a stream, use `fileno()`:

`#include <stdio.h>`

`int fileno(FILE *stream);`

- Upon success, `fileno()` returns the file descriptor associated with stream. On failure it returns -1

```
#include <stdio.h>
void main() {
    FILE *fp = fopen("file.txt", r);
    printf("File number is %d\n", fileno(fp));
    fclose(fp);
}
```

Controlling the Buffering

- Standard I/O implements three types of user buffering and provides developers with an interface for controlling the type and size of the buffer.
- The different types of user buffering serve different purposes:
 - Unbuffered : No buffering – characters are transmitted to the system as they are written
 - Line-buffered – characters are transmitted to the system as a block when a new-line character is encountered
 - Block-buffered: characters are transmitted to the system as a block when a buffer is filled.

Controlling buffering: setvbuf()

- setvbuf() function sets the buffering type of stream to mode:
#include <stdio.h>
int setvbuf(FILE *stream, char *buf, int mode, size_t size);
- Parameter
 - mode must be one of the following:
 - _IONBF : Unbuffered
 - _IOLBF : Line buffered
 - _IOFBF : Block buffered
 - buf : pointer to buffer (if NULL, buffer is automatically allocated by glibc)
 - size: buffer size in bytes

Controlling Buffering Errors

- Common Mistakes: supplied buffer must exist when stream is closed
 - Declaring buffer automatic variable in a scope that ends before stream is closed.
 - Be careful not to provide a buffer local to main and then fail to explicit close the stream.

```
#include <stdio.h>
void main() {
    char buf[BUFSIZ]; /* ERROR: local to main → global */
    /* set stdout to block-buffered with BUFSIZ buffer */
    setvbuf(stdout, buf, _IOFBF, BUFSIZ);
    printf("Arrr!Wn");
    return(0);
}
```

Thread Safety

- In computer programming, thread-safe describes a program portion or routine that can be called from multiple programming threads without unwanted interaction between the threads
- By using thread-safe routines, the risk that one thread will interfere and modify data elements of another thread is eliminated by circumventing potential data race situations with coordinated access to shared data.
- The standard I/O functions are inherently thread-safe
- Any given thread must acquire the lock and become the owning thread before issuing I/O requests
- Two or more threads operating on the same stream cannot interleave standard I/O operations, and thus within the context of single function calls, standard I/O operations are atomic.

Thread Safety: flockfile()

- Standard I/O provides a family of functions for individually manipulating the lock associated with a stream
 - Manual File Locking:
 - The function flockfile() waits until stream is no longer locked, bumps the lock count, and then acquires the lock, becoming the owning thread of the stream, and returns
- ```
#include <stdio.h>
void flockfile(FILE *stream);
```

# Thread Safety: funlockfile()

- funlockfile() decrements the lock count associated with stream:  
#include <stdio.h>  
void funlockfile(FILE \*stream);
- If lock count reaches zero, the current thread relinquishes ownership of the stream.
- Another thread is now able to acquire the lock

# Thread Safety: ftrylockfile()

- The ftrylockfile() function is a nonblocking version of flockfile()  
#include <stdio.h>  
**int ftrylockfile(FILE \*stream);**
- If stream is currently locked, ftrylockfile() does nothing and immediately returns a nonzero value
- If stream is not currently locked, it acquires the lock, bumps the lock count, becomes the owning thread of stream, and returns 0