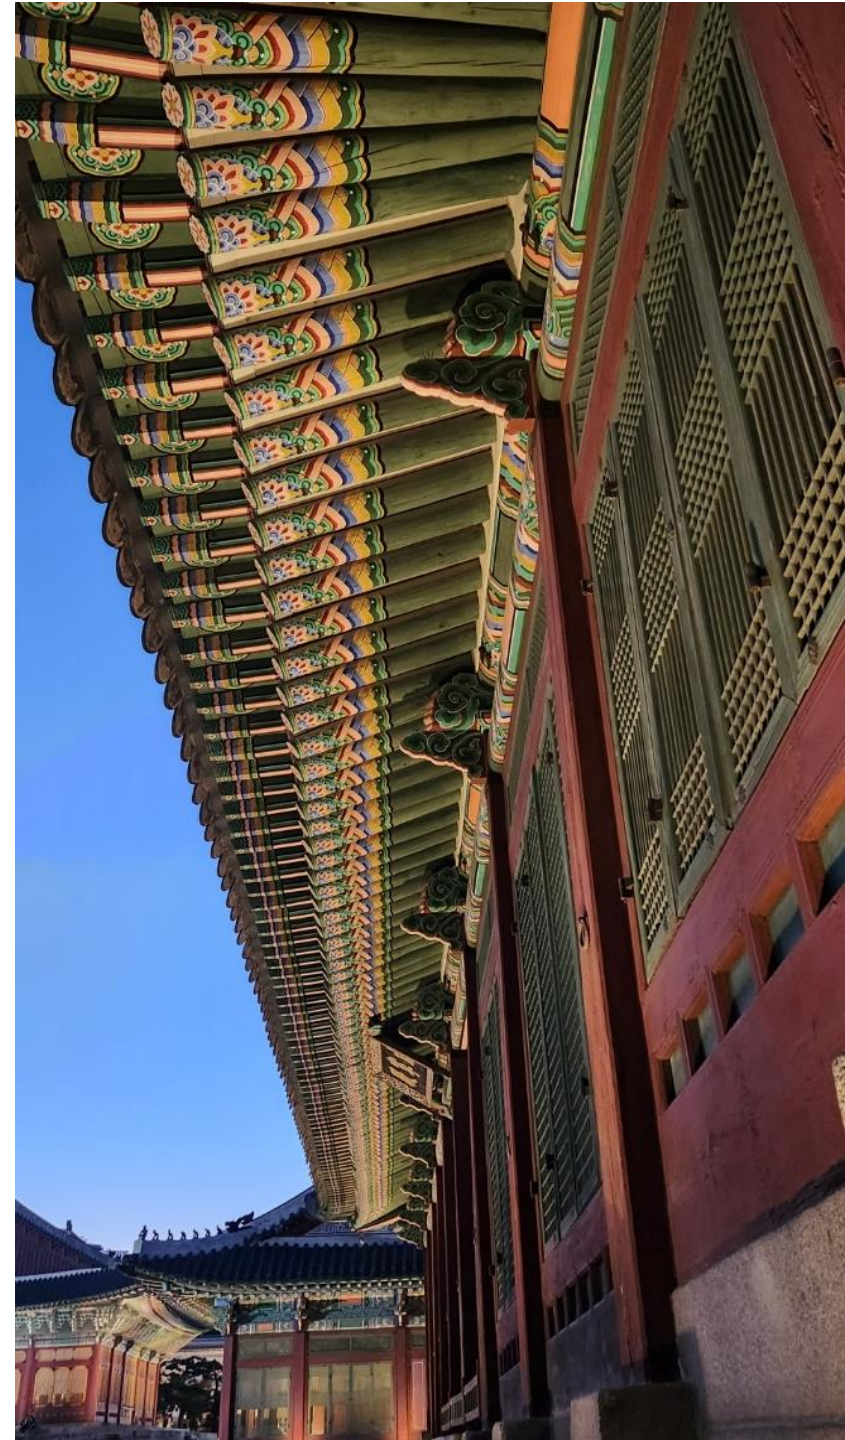


# Linux File System Programming

HGU



# File and filesystem in Linux

- File is the most fundamental and basic abstraction in Linux
- Linux follows 'Everything-is-a-file' philosophy (File, Directories, devices)
- Interaction often involve reading and writing file
- a File must first be opened for read or write.
- Open file are identified by a unique descriptor, file descriptor (fd)
- file descriptor is shared with user space, for file access

# Regular Files: file offset

- Regular Files are considered just as *Byte Stream*
  - no further organization or formatting is specified for a file. The bytes may have any values, and they may be organized within the file in any way.
- File Position
  - Read or Write Operation can be done in any Location within a file
  - The Location within a File is called the *File Position* or *File Offset*
  - When a file is first opened, the *File Position* is set zero.
  - As bytes are read from or written to the file, the file position increases accordingly, byte-by-byte.
  - The file position may also be set manually to a given value

# Regular File : length

- The size of a file is measured in bytes, and is called its *length*.
- The length is simply the number of bytes in the *linear array* that make up the file
- ***Length*** can be changed by '***truncation***' operation
  - A file can be truncated to a new size *smaller than its original size*, which results in bytes being *removed from the end of the file*.
  - A file can also be "*truncated*" to a new size *larger than its original size*. In that case, the new bytes are filled with zeros.
  - A file may be empty (have a *length of zero*).

# Regular File : Concurrent File Access

- A Single File can be opened more than once, by a different or even the same process.
- Each open instance of a File is given a unique file descriptor
- Processes can share their file descriptors, allowing a single descriptor to be used by more than one process
- Multiple Processes are free to read from and write to the same file at the same time. The results are generally *unpredictable*.

# i-node

- Instead of filename, file is referenced by i-node number (i-number) in Kernel
- An i-node stores metadata associated with a file:
  - modification timestamp,
  - owner,
  - type,
  - length,
  - the location of the file's data—but no filename!
- The i-node is both a physical object, located on disk in Unix-style filesystems, and a conceptual entity, represented by a data structure in the Linux kernel.

# Directory and Link

- **Accessing a file via its i-node number** is cumbersome (and also a potential security hole),
  - so files are always opened from user space *by a name*
- Directories are used to act as a *mapping* of human-readable names to i-node numbers.
  - A name and i-node pair is called a *link*.
  - Directory is a file that contains only a *mapping of names to i-nodes*
- When an application requests that a given filename be opened, the kernel opens the directory containing the filename and searches for the given name.
  - From the filename, the kernel obtains the i-node number.
  - From the i-node number, the i-node is found.
  - From the i-node, metadata of the file, including **on-disk location of the file** is found.
- Directories can be nested, creating a hierarchy, enabling pathnames like /home/user1/test

# The Stat Family

- a family of functions for obtaining the **metadata of a UNIX file**

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
#include <unistd.h>
```

```
int stat (const char *path, struct stat *buf); // returns information of file of path
```

```
int fstat (int fd, struct stat *buf); // returns information about the file of fd
```

```
int lstat (const char *path, struct stat *buf); // lstat() is identical to stat()
```

```
        // except in the case of symbolic link (link itself)
```



# struct stat

```
struct stat {  
    dev_t st_dev; /* ID of device containing file */  
    ino_t st_ino; /* file's inode number */  
    mode_t st_mode; /* file's mode bytes (permissions) */  
    nlink_t st_nlink; /* number of hard links at the file*/  
    uid_t st_uid; /* user ID of owner of the file */  
    gid_t st_gid; /* group ID of owner of the file */  
    dev_t st_rdev; /* device ID (if special file) */  
    off_t st_size; /* total size in bytes */  
    blksize_t st_blksize; /* blocksize for filesystem I/O */  
    blkcnt_t st_blocks; /* number of blocks allocated */  
    time_t st_atime; /* last access time */  
    time_t st_mtime; /* last modification time */  
    time_t st_ctime; /* last status change time */  
};
```

stat structure declared  
in <bits/stat.h>

# Sample Code (Retrieving the Size of a File)

```
#include <sys/types.h>
#include <sys/stat.h>
#include <unistd.h>
#include <stdio.h>

int main (int argc, char *argv[]) {
    struct stat sb;
    int ret;
    if (argc < 2) {
        fprintf(stderr, "usage: %s <file>\n", argv[0]);
        return 1;
    }
```

```
    ret = stat(argv[1], &sb);
    if(ret) {
        perror("stat");
        return 1;
    }
    printf("%s is %ld bytes\n", argv[1], sb.st_size);
    return 0;
}
```

```
yk@peace:~/systemprj/fileio_test$ ls -l stat1
-rwxrwxr-x 1 yk yk 8904 7월 30 11:35 stat1
yk@peace:~/systemprj/fileio_test$ ./stat1 stat1
stat1 is 8904 bytes
```

# Permissions

chmod() and fchmod() set a file's permissions to mode

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
int chmod(const char *path, mode_t mode);
```

```
int fchmod(int fd, mode_t mode);
```

```
// On success, all three calls return 0.
```

```
// On Failure, returns -1
```

# Sample Code to change file mode

```
#include <sys/types.h>
#include <sys/stat.h>
int ret;
main(int argc, char *argv[]) {
    /* set map.png in the current directory to
     * owner-readable and writable. (same as chmod 600 ./map.png)
     */
    ret = chmod("./map.png", S_IRUSR | S_IWUSR);
    if (ret) {
        perror("chmod");
        return 1;
    }
    return 0;
}
```

# Ownership

- change the field of `st_uid` and `st_gid` to change the owner and group of a file respectively.

```
#include <sys/types.h>
```

```
#include <unistd.h>
```

```
// changes the owner and group of a file
```

```
// if owner and group is -1, it is not changed by chown and lchown
```

```
// On success, all three calls return 0. On Failure, returns -1
```

```
int chown (const char *path, uid_t owner, gid_t group);
```

```
int lchown (const char *path, uid_t owner, gid_t group);
```

```
    // lchown changes owner of symbolic link itself not the original file
```

```
int fchown (int fd, uid_t owner, gid_t group);
```

# Sample Code to change Ownership

```
#include <stdio.h>
#include <unistd.h>
#include <grp.h>
int main( ) {
    struct group *gr;
    int ret;

    gr = getgrnam("prof");
    if(gr == NULL) {
        perror("getgrnam");
        return 1;
    }
    /* set mainifest.txt's group to 'prof */
    ret = chown ("chown-test.txt", -1, gr->gr_gid);
    if(ret != 0)
        perror("chmod");
    return 0;
}
```

```
yk@peace:~/systemprj/fileio_test$ ls -l
total 16
-rwxrwxr-x 1 yk yk 8712  7월 30 11:22 chown1
-rw-rw-r-- 1 yk yk 481  7월 30 11:22 chown1.c
-rw-rw-r-- 1 yk yk 0  7월 30 08:55 chown-test.txt
yk@peace:~/systemprj/fileio_test$ sudo ./chown1
[sudo] password for yk:
yk@peace:~/systemprj/fileio_test$ ls -l
total 16
-rwxrwxr-x 1 yk yk 8712  7월 30 11:22 chown1
-rw-rw-r-- 1 yk yk 481  7월 30 11:22 chown1.c
-rw-rw-r-- 1 yk prof 0  7월 30 08:55 chown-test.txt
yk@peace:~/systemprj/fileio_test$ |
```

# Directories

- Directory contains a list of filename, each of which maps to an inode number
  - Each name is called directory entry
  - Each name-to-inode mapping is called a link
- Directory contents are a listing of all the filenames in that directory
- Steps for Opening a file in a given directory
  1. Look up filename in that directory
  2. find corresponding inode number
  3. pass the inode number to the filesystem to find physical location of the file
- Directories can contain **other directories** (subdirectory)
  - All the directories (except /) are subdirectory of their parent directory
  - tow special subdirectories: dot(.) and dot-dot(..)
- pathname
  - absolute pathname : begins with the / directory
  - relative pathname : begins with **current working directory**

# Current Working Directory

- Def: **The starting point of the relative pathname**
- Every process has a **current working directory** which it initially inherits from its parent process.
- A process can both obtain and change the current working directory

```
#include <unistd.h>
char * getcwd(char *buf, size_t size);
```

- On success, copies the current working directory as an **absolute pathname** into buffer (buf) and returns pointer to buf
- On failure, returns NULL and set errno

```
// cwd.c
main() {
    char cwd[BUF_LEN];

    if (!getcwd (cwd, BUF_LEN)) {
        perror ("getcwd");
        exit (EXIT_FAILURE);
    }
    printf ("cwd = %s\n", cwd);
}
```



# Change the Current Working Directory

- When a user log into a system, the starting current working directory is set as the home directory (/etc/passwd)
- Linux provides two system calls for **changing current working directory**

```
#include <unistd.h>
```

```
int chdir (const char *path);
```

```
// changes current working directory by a given path
```

```
int fchdir (int fd); // changes cwd with file (fd)
```

- On success returns 0
- On failure returns -1

# Sample Code with Current Working Directory

```
char *swd;
int ret;

/* save the current working directory */
if(!(swd = getcwd (NULL, 0)){
    perror ("getcwd");
    exit (EXIT_FAILURE);
}

/* change to a different directory */
if(chdir (some_other_dir) <0) {
    perror ("chdir");
    exit (EXIT_FAILURE);
}

/* do some other work in the new directory...
*/
```

```
/* return to the saved directory */
if(!chdir (swd) ) {
    perror ("chdir");
    exit (EXIT_FAILURE);
}
free (swd)
```

# Create Directories

```
#include <sys/stat.h>
```

```
#include <sys/types.h>
```

```
int mkdir (const char *path, mode_t mode );
```

- mkdir() creates the directory **path**, which may be relative or absolute, with the permission bits **mode** (as modified by the current umask)
  - $\text{permission} = (\text{mode} \& \sim\text{umask} \& 01777)$
- Return Value
  - 0 on Success
  - -1 on Failure

# Remove Directories

include <unistd.h>

**int rmdir (const char \*path);**

- rmdir( ) removes path from the filesystem once its files have been removed.
- Return Value
  - 0 on Success
  - -1 on Failure

```
int ret;  
/* remove the directory /home/barbary/maps  
*/  
ret = rmdir ("/home/barbary/maps");  
if (ret)  
    perror ("rmdir");
```

# Read Directory's Contents

```
#include <sys/types.h>
```

```
#include <dirent.h>
```

```
DIR * opendir (const char *name);
```

- opendir() creates a directory stream representing the directory

```
#define _BSD_SOURCE /* or _SVID_SOURCE */
```

```
#include <sys/types.h>
```

```
#include <dirent.h>
```

```
int dirfd (DIR *dir);
```

- On success dirfd( ) returns the file descriptor backing the directory stream.
- On Failure it returns -1

# Reading from a Directory Stream

```
#include <sys/types.h>
```

```
#include <dirent.h>
```

```
struct dirent * readdir (DIR *dir);
```

- Once created directory stream with `opendir()`, program can begin reading entries from the directory.
- A successful call to `readdir( )` returns entry **one by one** from a given DIR. The `dirent` structure represents a directory entry.
- `readdir()` return NULL on End-of-List or Error (`errno=EBADF` : invalid dir)

```
// Linux definition for dirent:  
struct dirent {  
    ino_t d_ino; /* inode number */  
    off_t d_off; /* offset to the next dirent */  
    unsigned short d_reclen; /* length of this record */  
    unsigned char d_type; /* type of file */  
    char d_name[256]; /* filename */  
};
```

# Sample Code(1) for Reading Directory Stream

```
#include <stdio.h>
#include <dirent.h>
#include <sys/stat.h>
#include <string.h>
#include <stdlib.h>
#include <time.h>

void display_file_info(
    const char *path,
    const struct dirent *entry)
{
    struct stat st ; //file_stat
    char full_path[1024];

    // Create the full path to the file
    snprintf(full_path, sizeof(full_path),
        "%s/%s", path, entry->d_name);
```

```
drwxr-xr-x 3 4096 Sep 19 12:45 mydir
-rw-r--r-- 1 1234 Sep 18 10:30 myfile.txt
-rwxr-xr-x 1 4321 Sep 17 08:20 script.sh
```

```
// Get file statistics using stat()
if (stat(full_path, &st) == -1)
    perror("stat");
    return;
}

// Display file type and permission
printf("%c%c%c%c%c%c%c%c%c ",
    S_ISDIR(st.st_mode) ? 'd' : '-',
    st.st_mode & S_IRUSR ? 'r' : '-',
    st.st_mode & S_IWUSR ? 'w' : '-',
    st.st_mode & S_IXUSR ? 'x' : '-',
    st.st_mode & S_IRGRP ? 'r' : '-',
    st.st_mode & S_IWGRP ? 'w' : '-',
    st.st_mode & S_IXGRP ? 'x' : '-',
    st.st_mode & S_IROTH ? 'r' : '-',
    st.st_mode & S_IWOTH ? 'w' : '-',
    st.st_mode & S_IXOTH ? 'x' : '-');
```

```
// Display number of links
printf(" %lu", st.st_nlink);

// Display file size
printf(" %5ld", st.st_size);

// Display modification time
char time_buff[80];
struct tm *time_info;
time_info = localtime(&st.st_mtime);
strftime(time_buff, sizeof(time_buff),
    "%b %d %H:%M", time_info);

printf(" %s", time_buff);

// Display file name
printf(" %s\n", entry->d_name);
} // end of display_file_info
```

# Sample Code(2) for Reading Directory Stream

```
int main(int argc, char *argv[]) {  
    DIR *dir;  
    struct dirent *entry;  
  
    // Use current directory if no directory is specified  
    const char *dir_path = (argc > 1) ? argv[1] : ".";  
  
    // Open the directory stream  
    if ((dir = opendir(dir_path)) == NULL) {  
        perror("opendir");  
        return EXIT_FAILURE;  
    }  
}
```

```
    // Read directory entries and display them one by one  
    while ((entry = readdir(dir)) != NULL) {  
        if (strcmp(entry->d_name, ".") != 0 &&  
            strcmp(entry->d_name, "..") != 0) {  
            display_file_info(dir_path, entry);  
        }  
    }  
    // Close the directory stream  
    closedir(dir);  
    return EXIT_SUCCESS;  
}
```



# Closing the Directory Stream

```
#include <sys/types.h>
```

```
#include <dirent.h>
```

```
int closedir (DIR *dir);
```

- `closedir()` closes directory `dir` that was open by `diropen()`

# Sample Code for Reading Directory Stream

```
int find_file_in_dir (const char *path, const char *file)
{
    struct dirent *entry;
    int ret = 1;
    DIR *dir;

    dir = opendir (path);
    errno = 0;
    while ((entry = readdir (dir)) != NULL) {
        if (!strcmp(entry->d_name, file)) {
            ret = 0;
            break;
        }
    }
    if (errno && !entry)
        perror ("readdir");
    closedir (dir);
    return ret; // 0 on Success, non-Zero on Failure
}
```

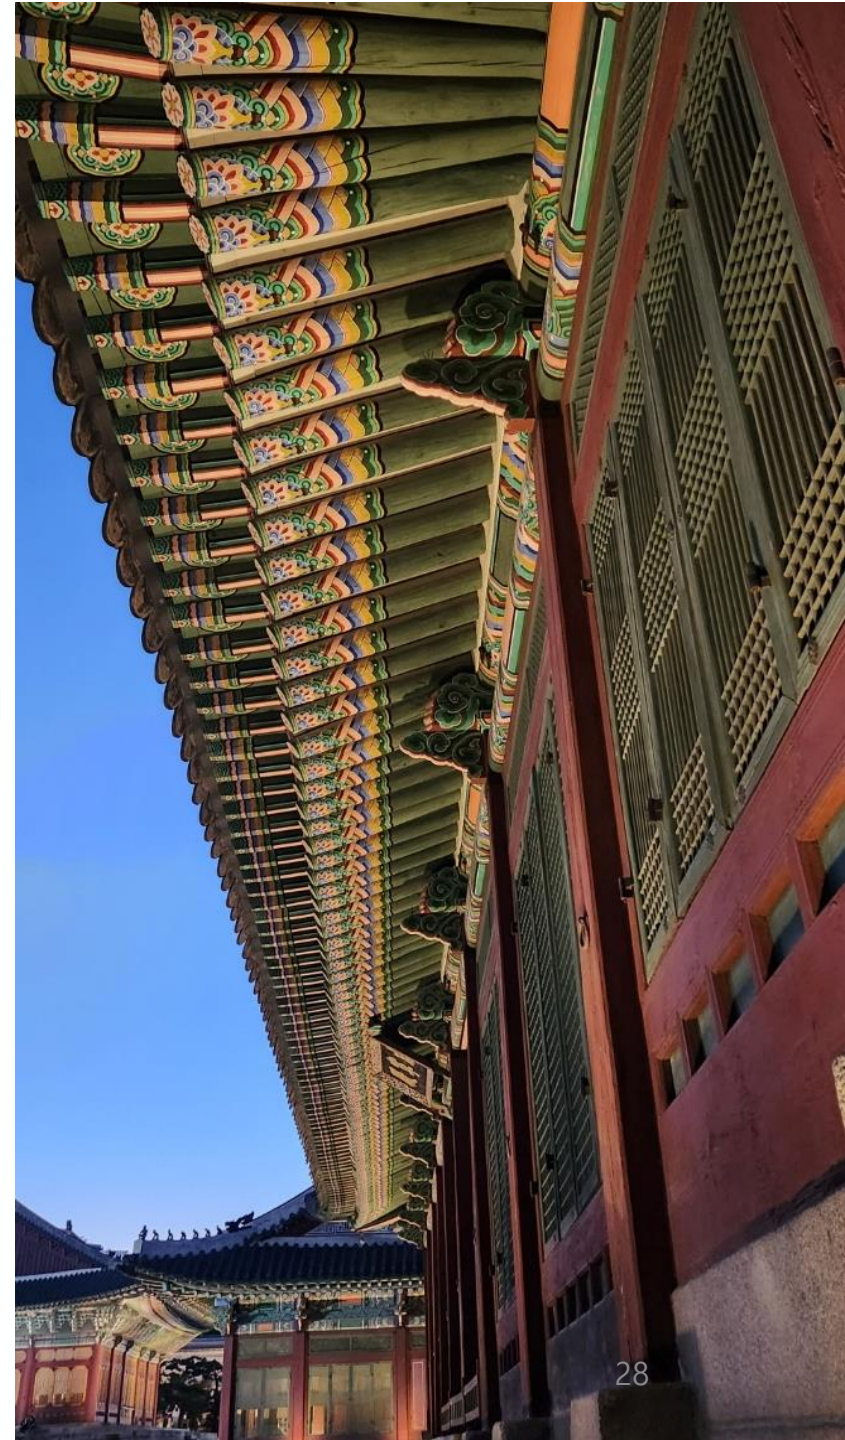
- find\_file\_in\_dir
  - searches the directory 'path' for a file
  - Returns 0 if 'file' exists in 'path' and a nonzero value otherwise

# Filesystem API Summary

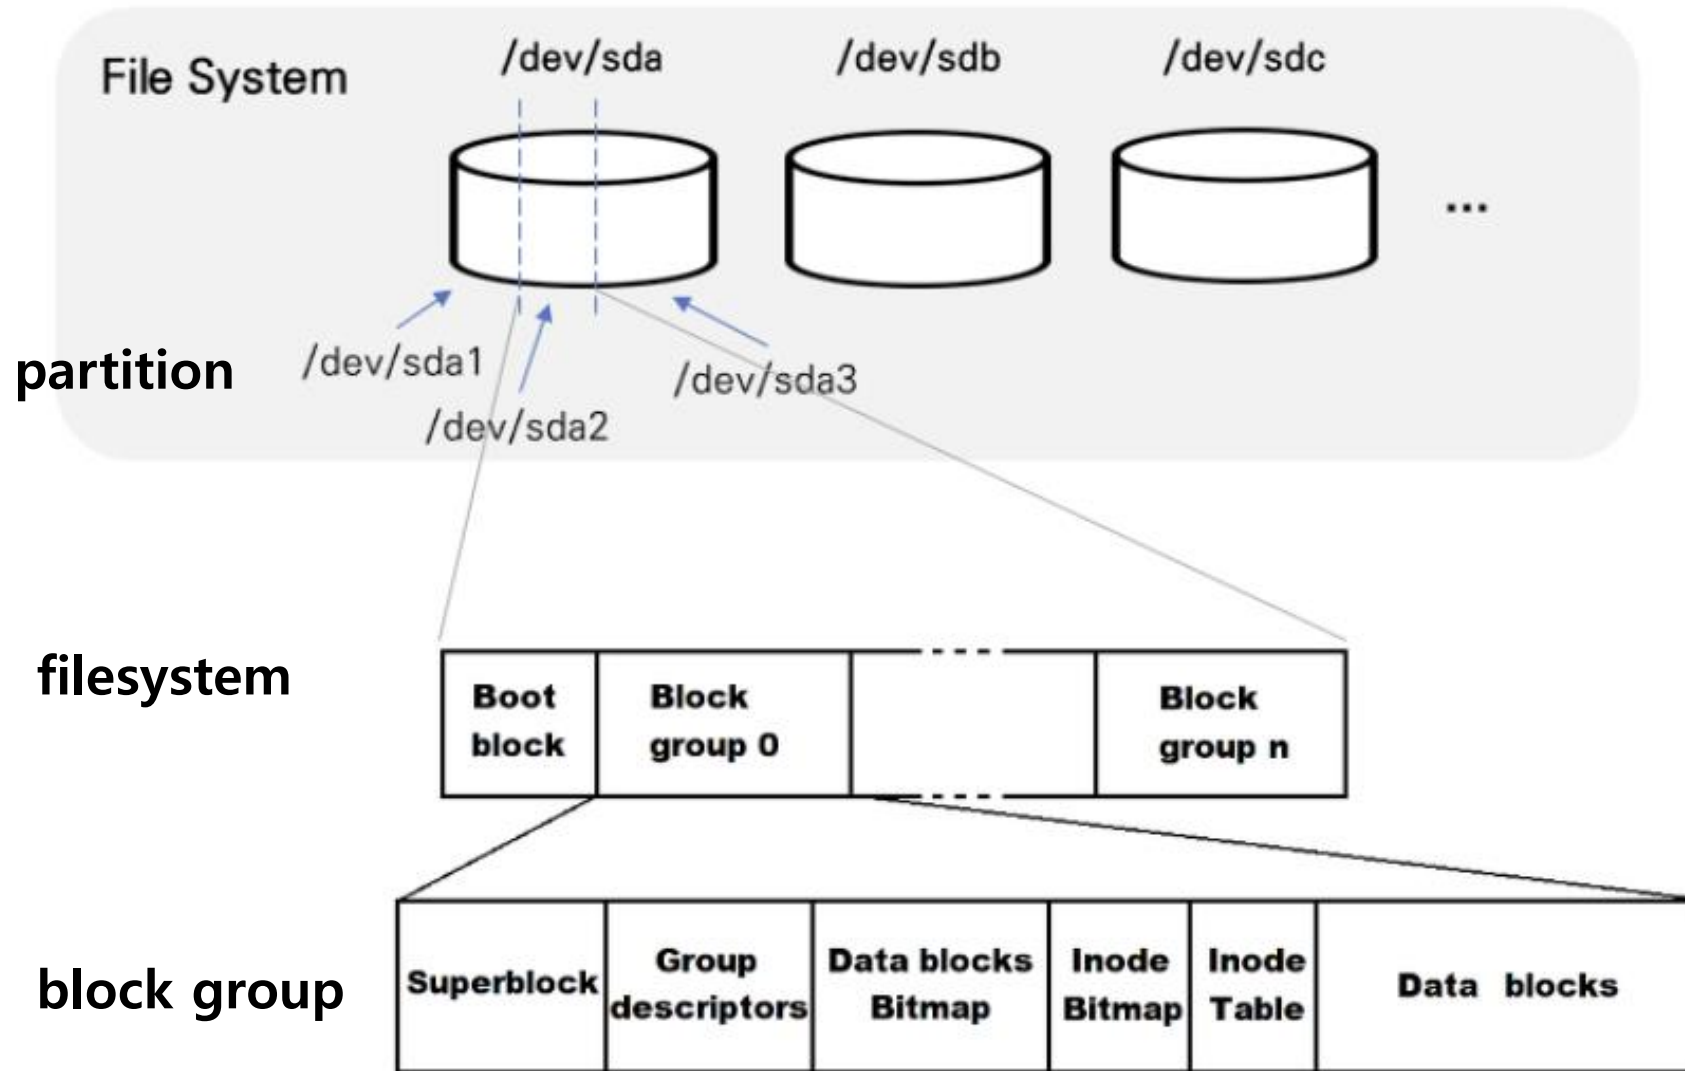
- `int stat (const char *path, struct stat *buf)`
- `int fstat (int fd, struct stat *buf)`
- `int chmod(const char *path, mode_t mode)`
- `int fchmod(int fd, mode_t mode)`
- `int chown (const char *path, uid_t owner, gid_t group)`
- `int fchown (int fd, uid_t owner, gid_t group)`
- `char * getcwd(char *buf, size_t size)`
- `int chdir (const char *path)`
- `int mkdir (const char *path, mode_t mode )`
- `int rmdir (const char *path)`
- `DIR * opendir (const char *name)`
- `int closedir (DIR *dir)`
- `struct dirent * readdir (DIR *dir)`

# Linux File System Internals

HGU



# File System: Disk and Partition



- Disk
- Partition
- Boot block
- Block Groups
  - super block
  - group descriptors
  - datablock bitmap
  - inode bitmap
  - inode table
  - data block

# Making Filesystem (Linux Commands)

- Add New disk (SCSI disk)
  - `echo "- - -" > /sys/class/scsi_host//host0/scan`
- Making Partition
  - `fdisk -l` : list disk partition states
  - **`fdisk /dev/sdb`**
- Device should be formatted
  - **`mkfs /dev/sdb1`**
- Each Device should be mounted (linked) to a specific directory
  - **`mount /dev/sdb1 / disk1`**
  - `.etc/fstab` : automounting device list when booting system
- Checking Filesystem
  - `df -h`
- Mounted Device is accessed through a path from root (/)
  - `cd /disk1`

# What is i-node?

- Linux filesystem has 'blocks' which hold data called 'inode' (**index node**) describing **metadata** of file: ownership, permission, time, size, and location of file's data
- **i-node** means both
  - *physical block* located on disk in UNIX-like filesystem and
  - *conceptual entry* represented by data structure in Linux Kernel
- i-node is a **unique identifier for each file** (or directory) in a filesystem.



# i-node and file copy and move

- Checking i-node information by Linux commands:

```
yk@peace:~/systemprj/fileio_test$ touch mytest
yk@peace:~/systemprj/fileio_test$ stat mytest
  File: 'mytest'
  Size: 0                Blocks: 0          IO Block: 4096   regular empty file
Device: 801h/2049d      Inode: 21977514   Links: 1
Access: (0664/-rw-rw-r--)  Uid: ( 1017/   yk)   Gid: ( 1019/   yk)
Access: 2024-08-03 20:29:49.956590670 +0900
Modify: 2024-08-03 20:29:49.956590670 +0900
Change: 2024-08-03 20:29:49.956590670 +0900
Birth: -
yk@peace:~/systemprj/fileio_test$ ls -li mytest
21977514 -rw-rw-r-- 1 yk yk 0  8월  3 20:29 mytest
```

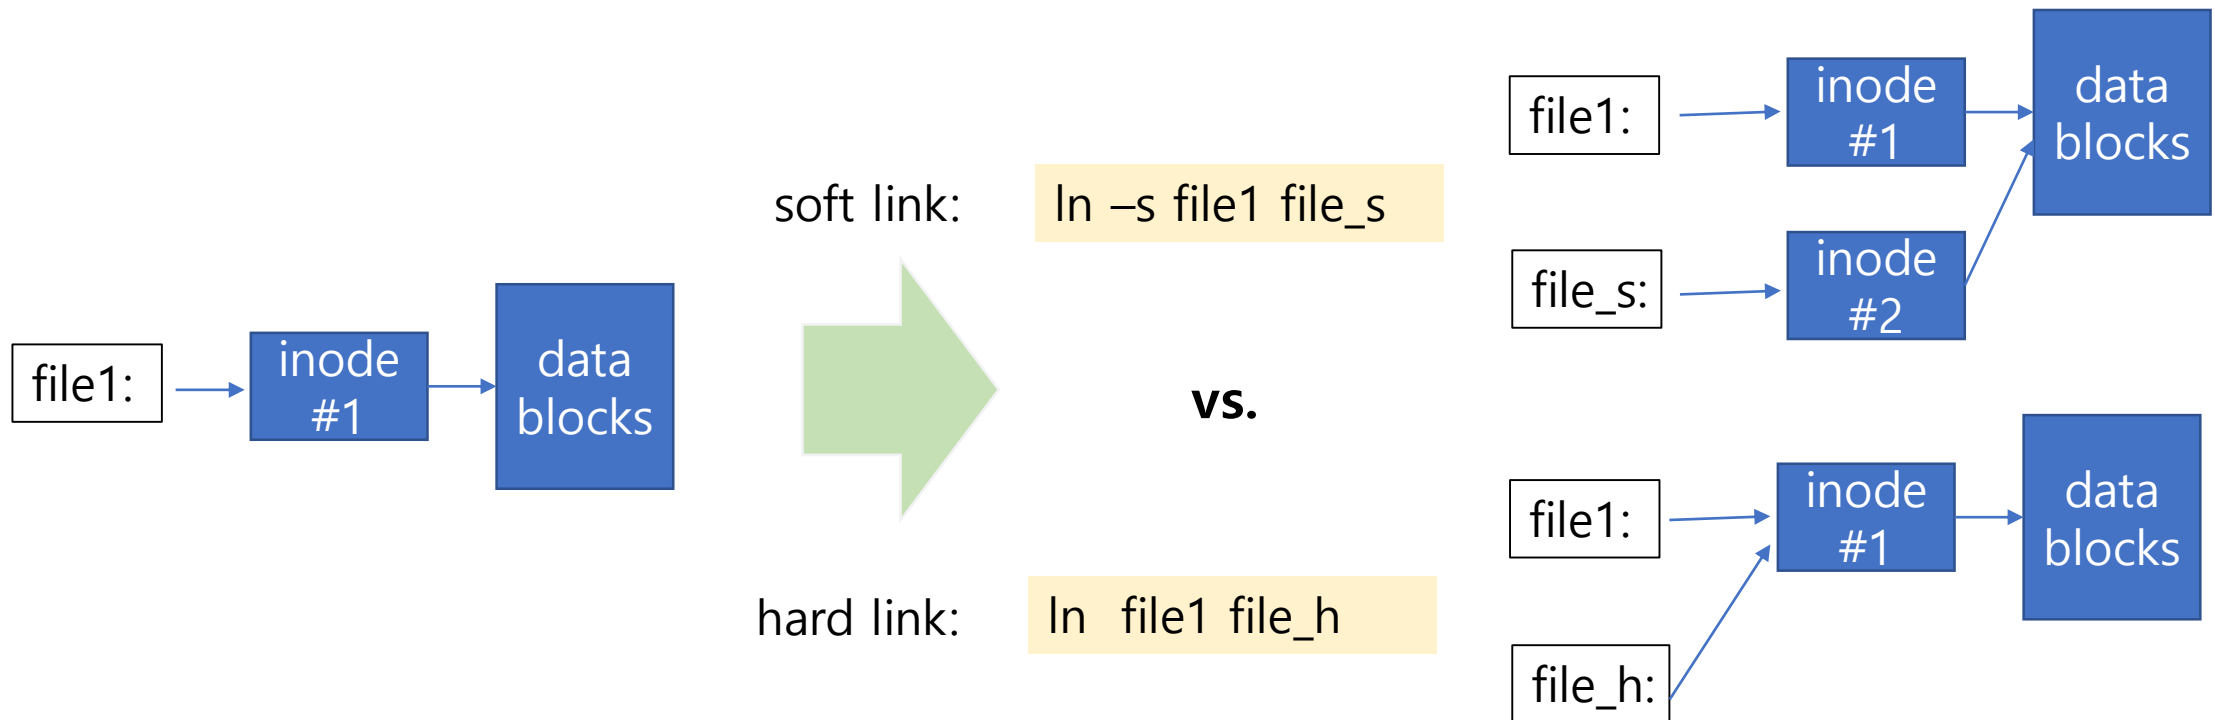
- i-node with cp and mv:

```
yk@peace:~/systemprj/fileio_test$ ls -li mytest
21977514 mytest
yk@peace:~/systemprj/fileio_test$ cp mytest mytest_cp
yk@peace:~/systemprj/fileio_test$ mv mytest mytest_mv
yk@peace:~/systemprj/fileio_test$ ls -li mytest*
21977512 -rw-rw-r-- 1 yk yk 0  8월  6 07:37 mytest_cp
21977514 -rw-rw-r-- 1 yk yk 0  8월  3 20:29 mytest_mv
```



# Link Commands and inode

- **soft link (symbolic link)**: create new i-node, points to same data blocks
- **hard link** : share i-node with different name, increase the reference count of inode



# Files in Linux

- Byte String
- Arbitrarily addressable
- contents has no predefined properties
- protected by access rights
  - r
  - w
  - x
- defined for user, group, others

# I-node (Index Node)

- Each file is represented by an inode
- contains
  - owner (UID, GID)
  - Access rights
  - Time of last accessed/modified
  - Size
  - Type (directory, file, device, pipe, ...)
  - pointers to data blocks that stores file's content
  - No file name!

# Directories

- Directories are handled as normal files, but are marked in inode-type as directory
- Directory entry contains
  - Length of the entry
  - Name (variable length up to 255 characters)
  - Inode number
- Multiple directory entries may reference the same Inode number (hard link)
- Users **identify files** via pathnames (".path/to/file") that are mapped to Inode numbers by OS
- If the path starts with "/", it is absolute and is resolved up from the root
- Otherwise the path is resolved relative to the current directory

# Directories

- Each directory contains an entry "." that represents the Inode of the current directory
- The second entry ".." references parent directory
- The path is resolved from left to right and the respective name is looked up in the directory

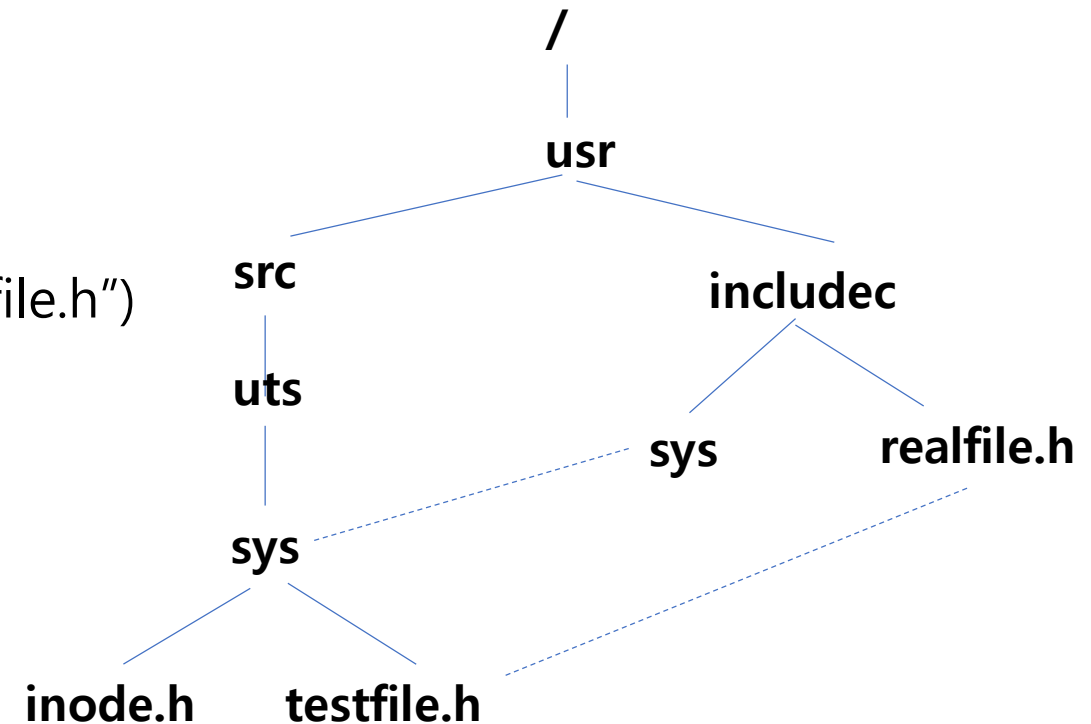
# Symbolic Links (softlink)

- To improve shared access to files, UNIX allows use of symbolic links to reference single files and directories via multiple different paths
- `symlink(existing_name, new_name)` creates an additional path to the resource

- Example:

`symlink("usr/src/uts/sys", "/usr/include/sys")`

`symlink("/usr/include/realfile.h", "/usr/src/uts/sys/testfile.h")`

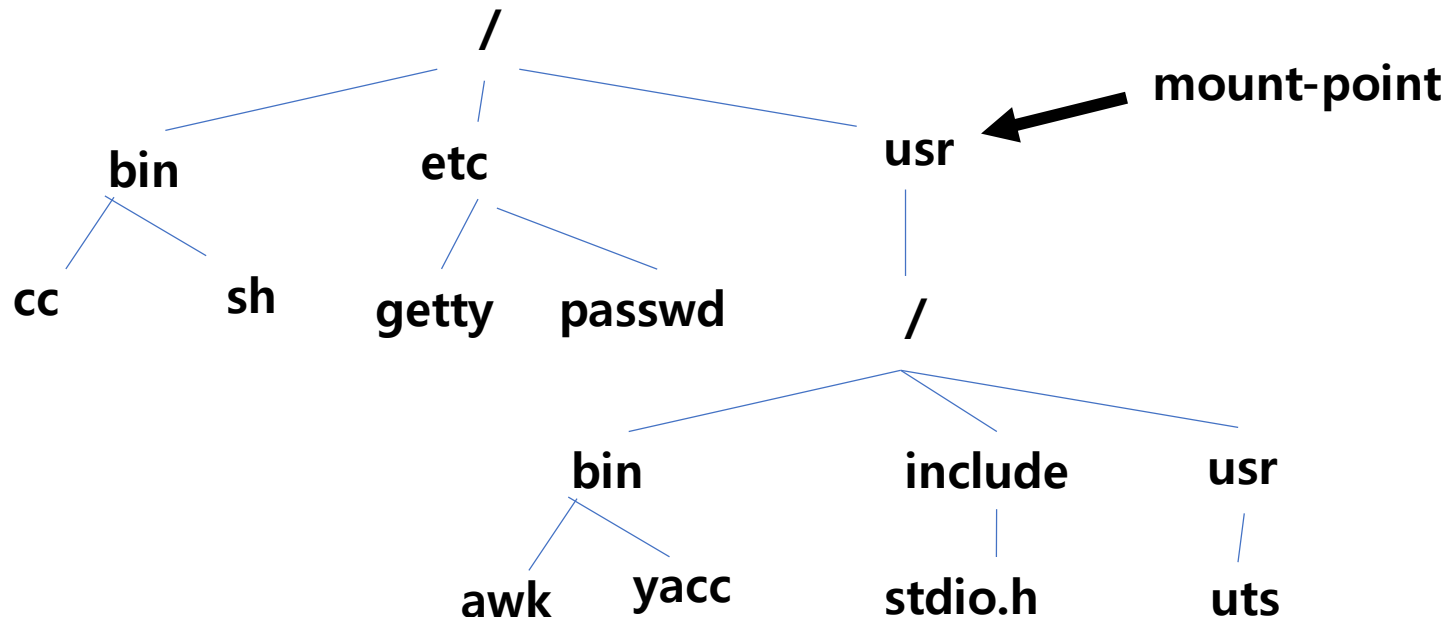


# Hard and Symbolic Links

- A hard link is an additional file name
  - there exists another directory entry that points to the same file
  - All hard links point to the same inode
  - Each new hard link increments the link count of the Inode
  - As long as the count  $\neq 0$ , the file "survives" a `remove()` and only link count is decremented
  - If last link is removed, the file is deleted and the Inode can be reused
- A Symbolic link (soft link) is a file that contains the path of another file/directory
  - Symbolic links are interpreted and resolved on every access
  - If target of a symbolic link is deleted, the link becomes invalid but remains existent

# Logical and Physical System

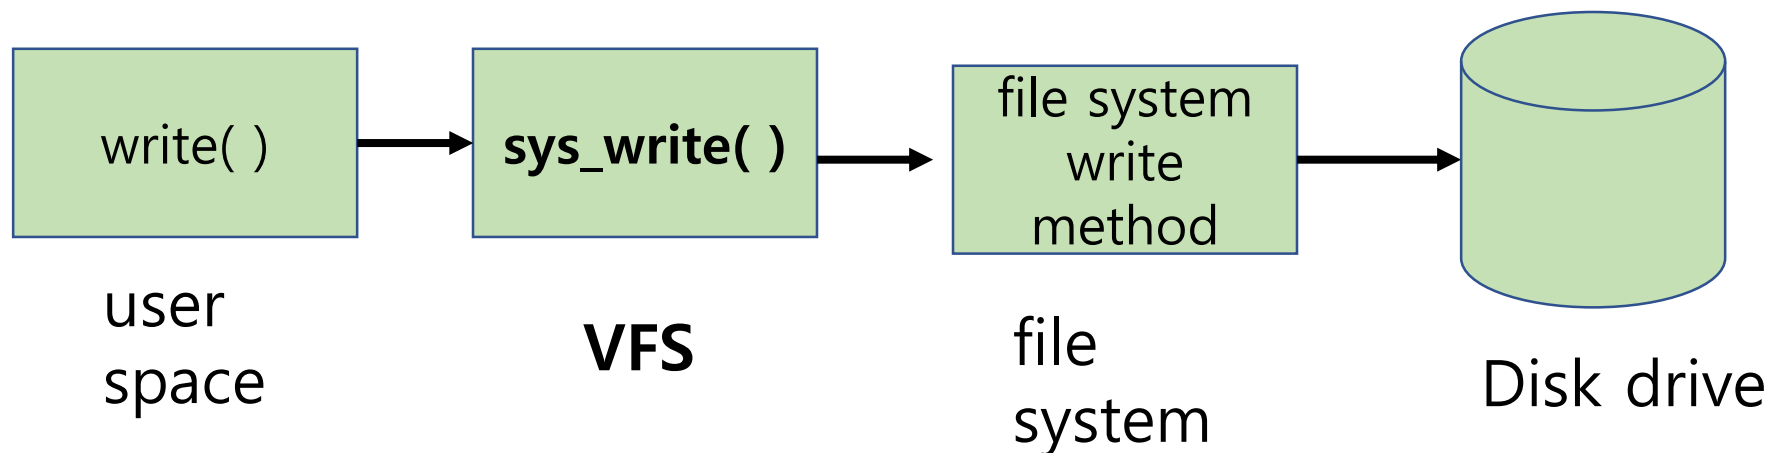
- A logical File system may consists of multiple physical file systems
- A file system can be hooked into any path of the virtual file system tree with the "mount" command
- Mounted file systems are managed by the OS in a "mount table" that connects path to mount points
- This allows to identify the root Inode of mounted file systems





# Virtual File System

- The **Virtual File System (VFS)** implements a generic file system interface between the actual file system implementation (in kernel) and accessing applications to provide interoperability
  - ➔ Applications can access different file systems on different media via a homogeneous set of UNIX system calls
- Example `write(f, &buf, len);`
  - translated into system call
  - system call is forwarded to the actual file system implementation



# VFS Objects and Data Structure

- VFS is object oriented
- Four base objects
  - Super block: represents specific properties of a filesystem
  - Inode: File description
  - Dentry: the directory entry represents a single component of a path
  - File: representation of an open file that is associated with a process
- VFS handles directories like files
  - Dentry object represents component of a path that may be a file
  - Directories are handled like files as Inode
- Each object provides a set of operations

# Superblock

- Each file system must provide a superblock
  - Contains properties of the file system
  - Is stored on special sectors of disk or is created dynamically
  - Structure is created by `alloc_super()` when the file system is mounted

```
struct super_block {
    struct list_head    s_list;           /* Keep this first */
    dev_t               s_dev;           /* search index; _not_
    unsigned long       s_blocksize;
    unsigned char       s_blocksize_bits;
    unsigned char       s_dirt;
    unsigned long long  s_maxbytes;      /* Max file size */
    struct file_system_type *s_type;
    struct super_operations *s_op;
    struct dquot_operations *dq_op;
    struct quotactl_ops  *s_qcop;
    struct export_operations *s_export_op;
    unsigned long       s_flags;
    unsigned long       s_magic;
    struct dentry        *s_root;
    struct rw_semaphore  s_umount;
    struct mutex         s_lock;
    int                 s_count;
    int                 s_syncing;
    int                 s_need_sync_fs;
    atomic_t            s_active;
    void                *s_security;
    struct xattr_handler **s_xattr;

    struct list_head    s_inodes;        /* all inodes */
    struct list_head    s_dirty;        /* dirty inodes */
    struct list_head    s_io;           /* parked for writeback
    struct hlist_head   s_anon;         /* anonymous dentries
    struct list_head    s_files;

    struct block_device *s_bdev;
    struct list_head    s_instances;
    struct dquot_info    s_dquot;      /* Diskquota specific
```

# I-node Object

- Contains information specific to a file
- For typical UNIX file systems, an inode can directly be read from disk
- Special Entries for non-data files
  - `i_pipe`, `i_bdev`, `i_cdev` are reserved for pipe, block device, charac device
- Some entries are not supported by all file systems and may therefore be set to Null

```
struct inode {  
    struct hlist_node    i_hash;  
    struct list_head     i_list;  
    struct list_head     i_sb_list;  
    struct list_head     i_dentry;  
    unsigned long         i_ino;  
    atomic_t             i_count;  
    umode_t              i_mode;  
    unsigned int          i_nlink;  
    uid_t                i_uid;  
    gid_t                i_gid;  
    dev_t                i_rdev;  
    loff_t               i_size;  
    struct timespec       i_atime;  
    struct timespec       i_mtime;  
    struct timespec       i_ctime;  
    unsigned int          i_blkbits;  
    unsigned long         i_blksize;  
    unsigned long         i_version;  
    unsigned long         i_blocks;  
    unsigned short        i_bytes;  
    spinlock_t            i_lock;  
    struct mutex          i_mutex;  
    struct rw_semaphore   i_alloc_sem;  
    struct inode_operations *i_op;  
    struct file_operations *i_fop;  
    struct super_block     *i_sb;  
    struct file_lock      *i_flock;  
};
```

# Dentry Object

- UNIX directories are handled like files
- The path /bin/vi contains the directory / and bin as well as file vi
- resolution of path requires introduction of dentry objects
- Each part of path is dentry object
- VFS creates dentry objects on the fly
- No equivalent on disk drive

```
struct dentry {
    atomic_t d_count;
    unsigned int d_flags;           /* protected
    spinlock_t d_lock;             /* per dentry
    struct inode *d_inode;         /* Where the
                                    * negative *

    /*
     * The next three fields are touched by __d_l
     * so they all fit in a cache line.
     */
    struct hlist_node d_hash;       /* lookup has
    struct dentry *d_parent;       /* parent dir
    struct qstr d_name;

    struct list_head d_lru;        /* LRU list *
    /*
     * d_child and d_rcu can share memory
     */
    union {
        struct list_head d_child;  /* ch
        struct rcu_head d_rcu;

    } d_u;
    struct list_head d_subdirs;    /* our childr
    struct list_head d_alias;     /* inode alia
    unsigned long d_time;         /* used by d_
    struct dentry_operations *d_op;
    struct super_block *d_sb;     /* The root o
```

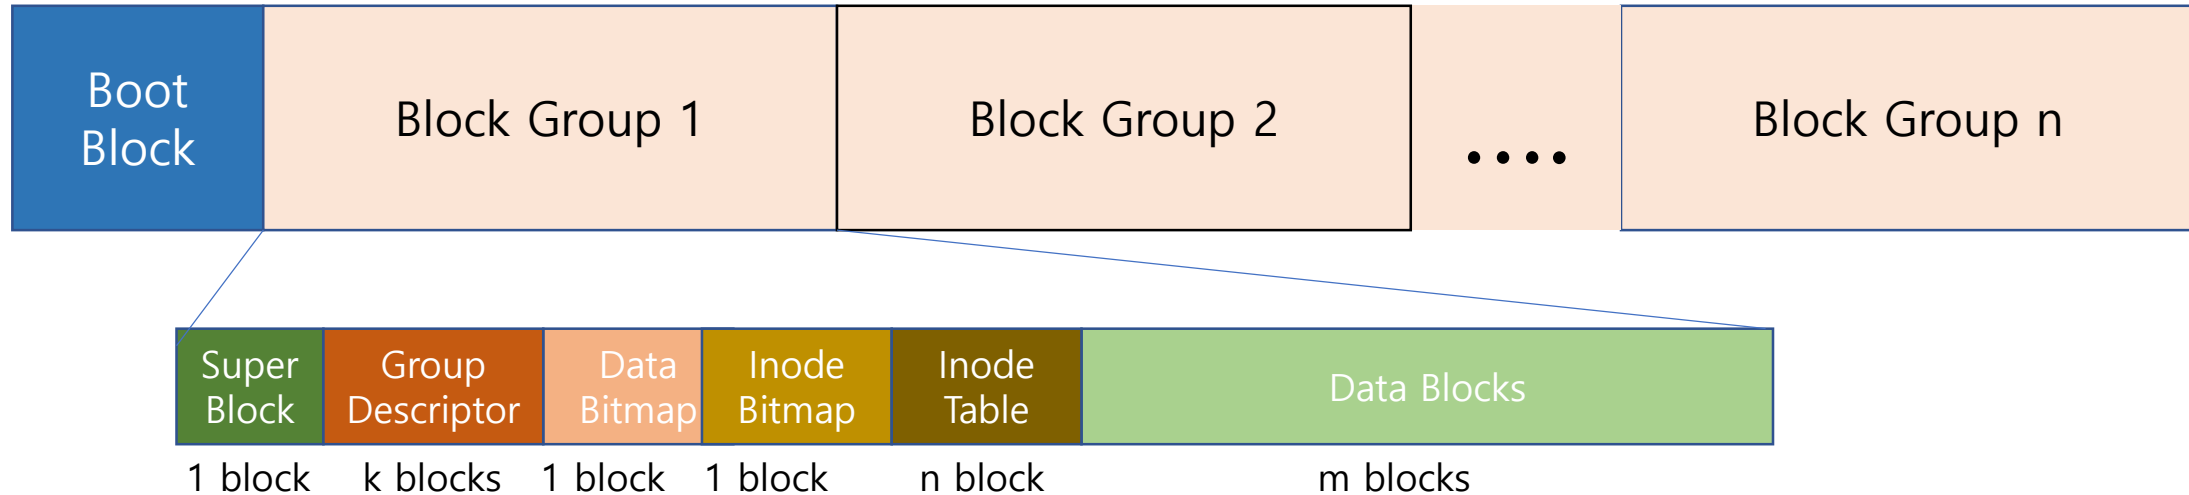
# File Object

- File object represents open file
- Interface to applications
- is created as reply to open() system call
- is removed on close() system call
- different processes can open a file multiple times → different file object
- The file object is an in-memory data structure of the OS

```
struct file {  
    union {  
        struct list_head      fu_list;  
        struct rcu_head       fu_rcuhead;  
    } f_u;  
    struct dentry              *f_dentry;  
    struct vfsmount            *f_vfsmnt;  
    struct file_operations     *f_op;  
    atomic_t                   f_count;  
    unsigned int               f_flags;  
    mode_t                     f_mode;  
    loff_t                     f_pos;  
    struct fown_struct          f_owner;  
    unsigned int               f_uid, f_gid;  
    struct file_ra_state       f_ra;  
    unsigned long              f_version;  
    void                       *f_security;  
    void                       *private_data;  
    struct list_head           f_ep_links;  
    spinlock_t                 f_ep_lock;  
    struct address_space        *f_mapping;  
};
```

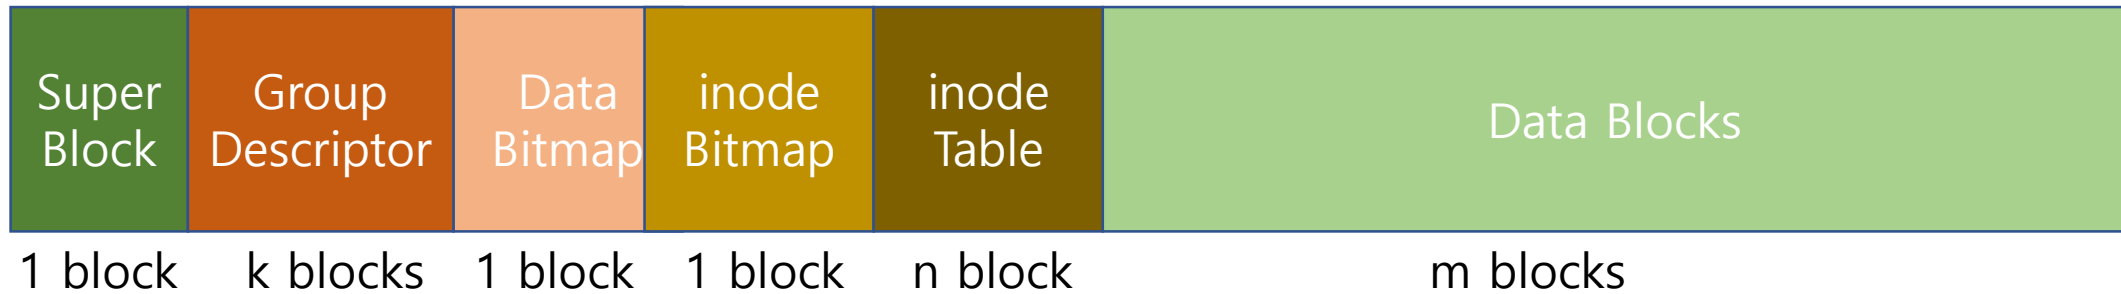
# EXT2 Architecture

- EXT2 divides storage system into block groups



- Boot block is equivalent to first sector on Hard Disk.
- Block Group is basic component, which contains further filesystem components

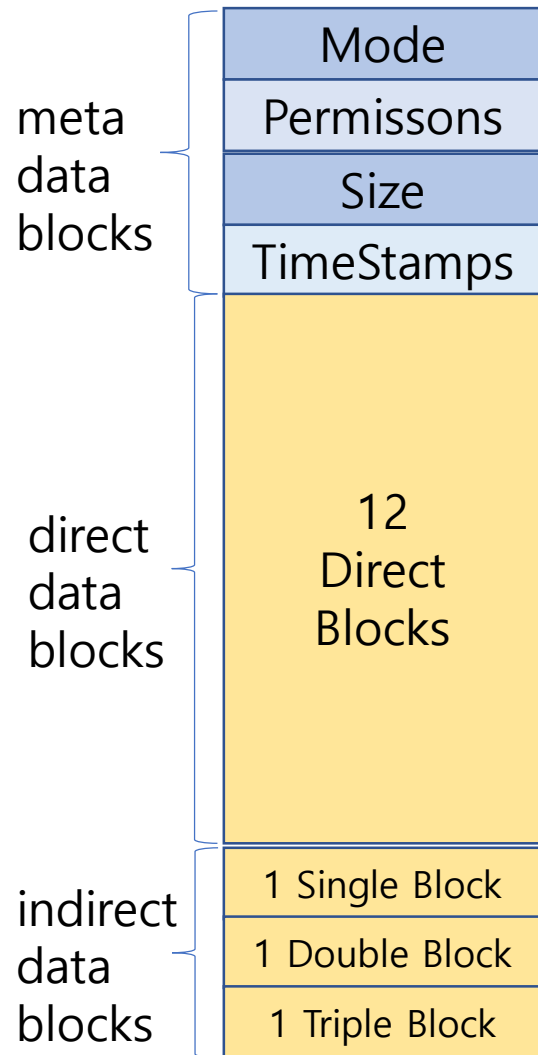
# Metadata for each Block Group



- Super Block: Central Structure, which contains number of free, allocated blocks, state of the filesystem, used block size, ...
- Group Descriptor: contains the state, number of free blocks and inodes in each block group. Each block group has group descriptor
- Data Bitmap: 1/0 allocation representation for data block
- inode Bitmap: 1/0 allocation representation for inode blocks
- inode Table: stores all inodes for this block group
- Data Blocks: User data, File Contents

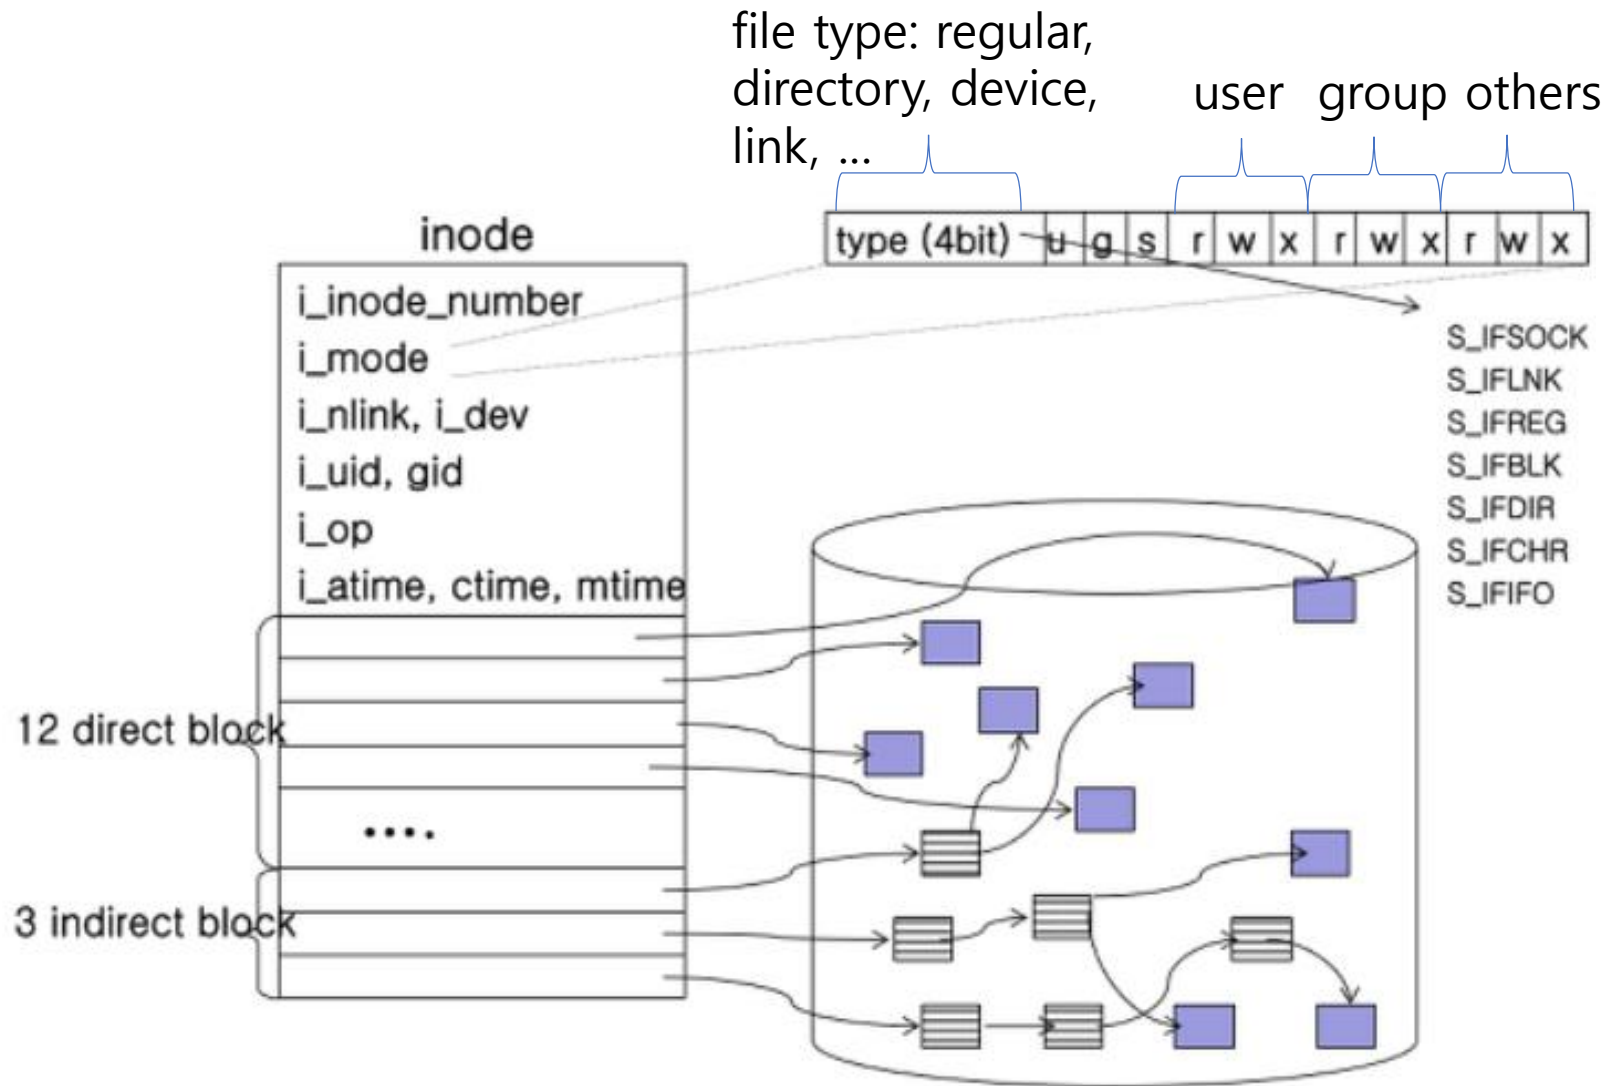


# i-node Structure



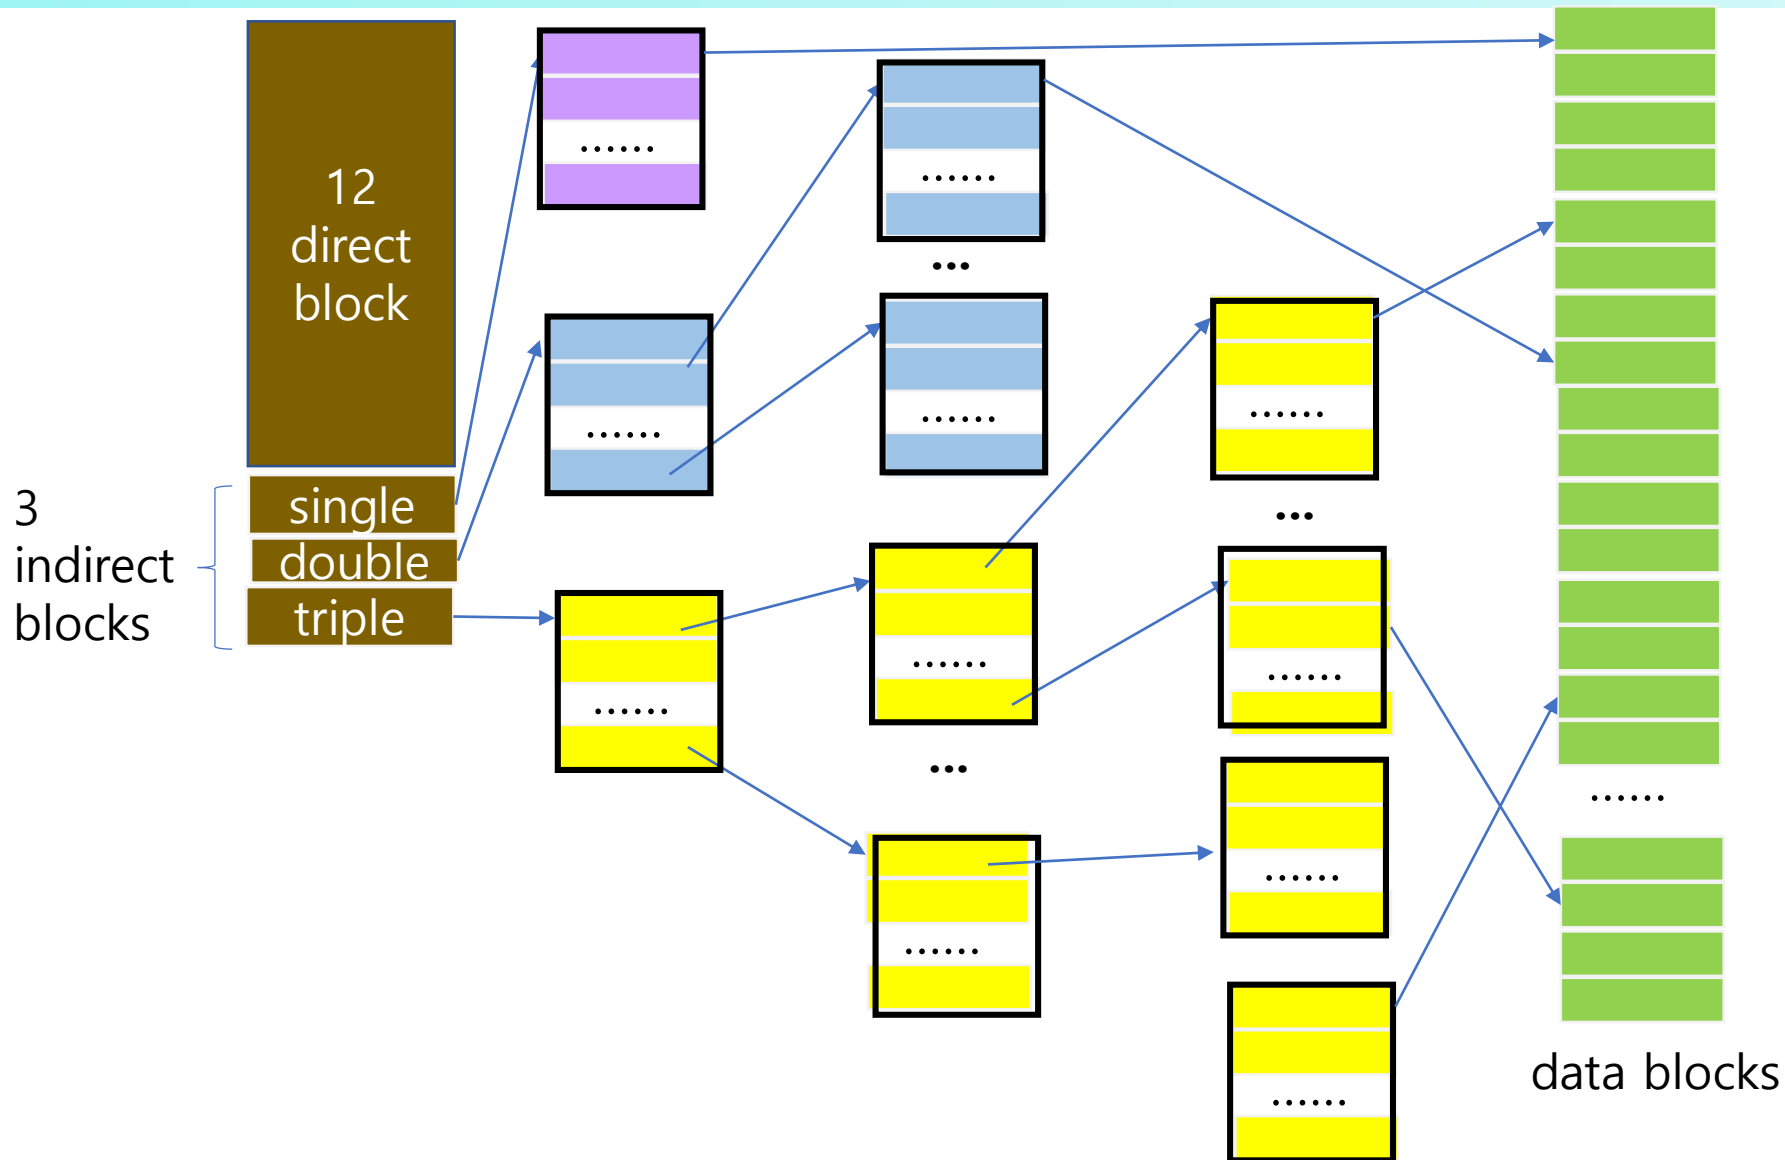
- mode : ownership of user, group
  - Permission: read, write, execute for each ownership
  - Size: file size in bytes
  - Time Stamps : creation, modification, recent access
  - Type: file, directory, device, pipe,...
- 
- Direct Blocks: pointer to file data block
  - Indirect Blocks: Single, double, triple blocks to file data block

# i-node metadata members



- i\_mode (16bits): file type
- i\_nlink: #of hard links
- i\_uid, gid: user id, group id
- i\_op : file operation functions
- i\_atime, c\_time, mtime: time of access, creation, modification
- i\_size: file size in bytes
- i\_blocks: #of data blocks allocated
- i\_block[15] : pointers to data blocks

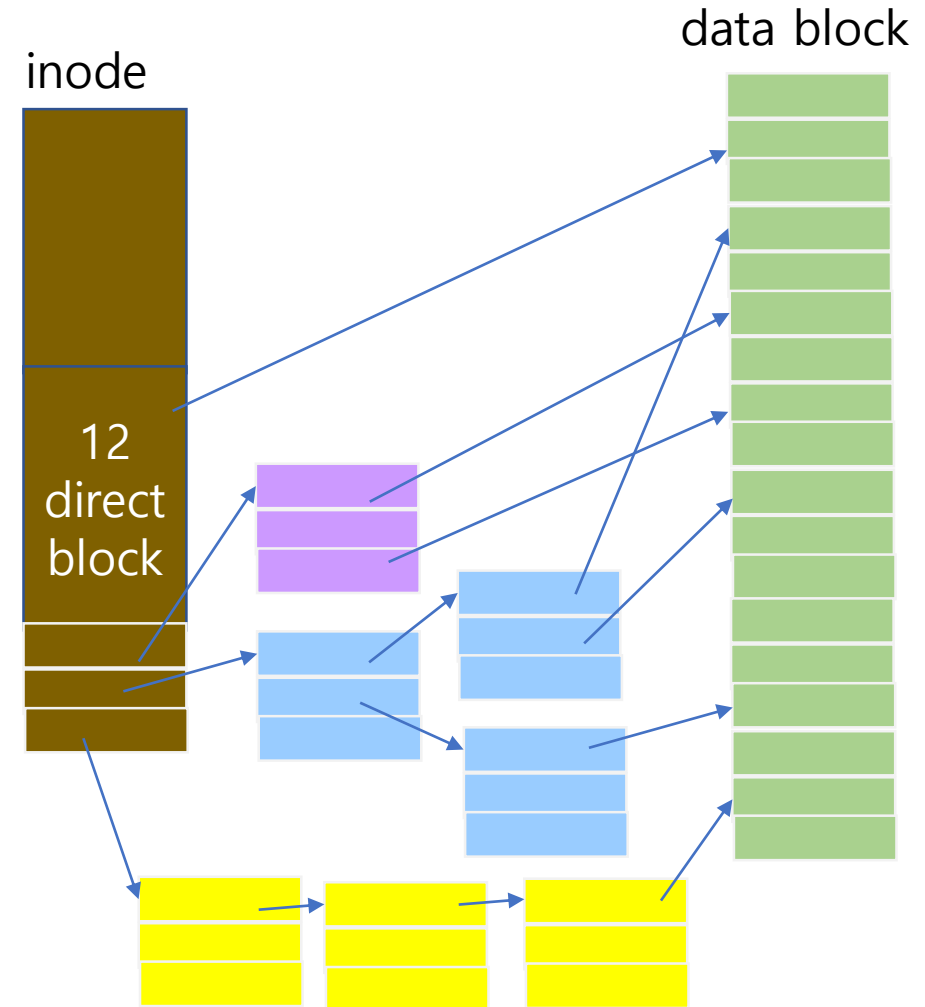
# i-node Data blocks



- 12 direct pointers to data block
- 3 indirect pointers
  - 1 single indirect ptrs
  - 1 double indirect ptrs
  - 1 triple indirect ptrs

# Limitation in File Size

- Size of a file is restricted by the number of block entries in inode
- Assumption:
  - 700 Mbyte file size and 4KB block size
  - 179,200 block entries ( $700 \times 1024 / 4$ ) are necessary and each entry needs 32bits(4 bytes/entry) → 700KB
  - If the Inode size is fixed, then you also need 700 KByte inode for 4Kbytes files.
- EXT2 filesystem supports direct and indirect blocks
  - There is one pointer each for one-time, two-time, and three time indirect blocks

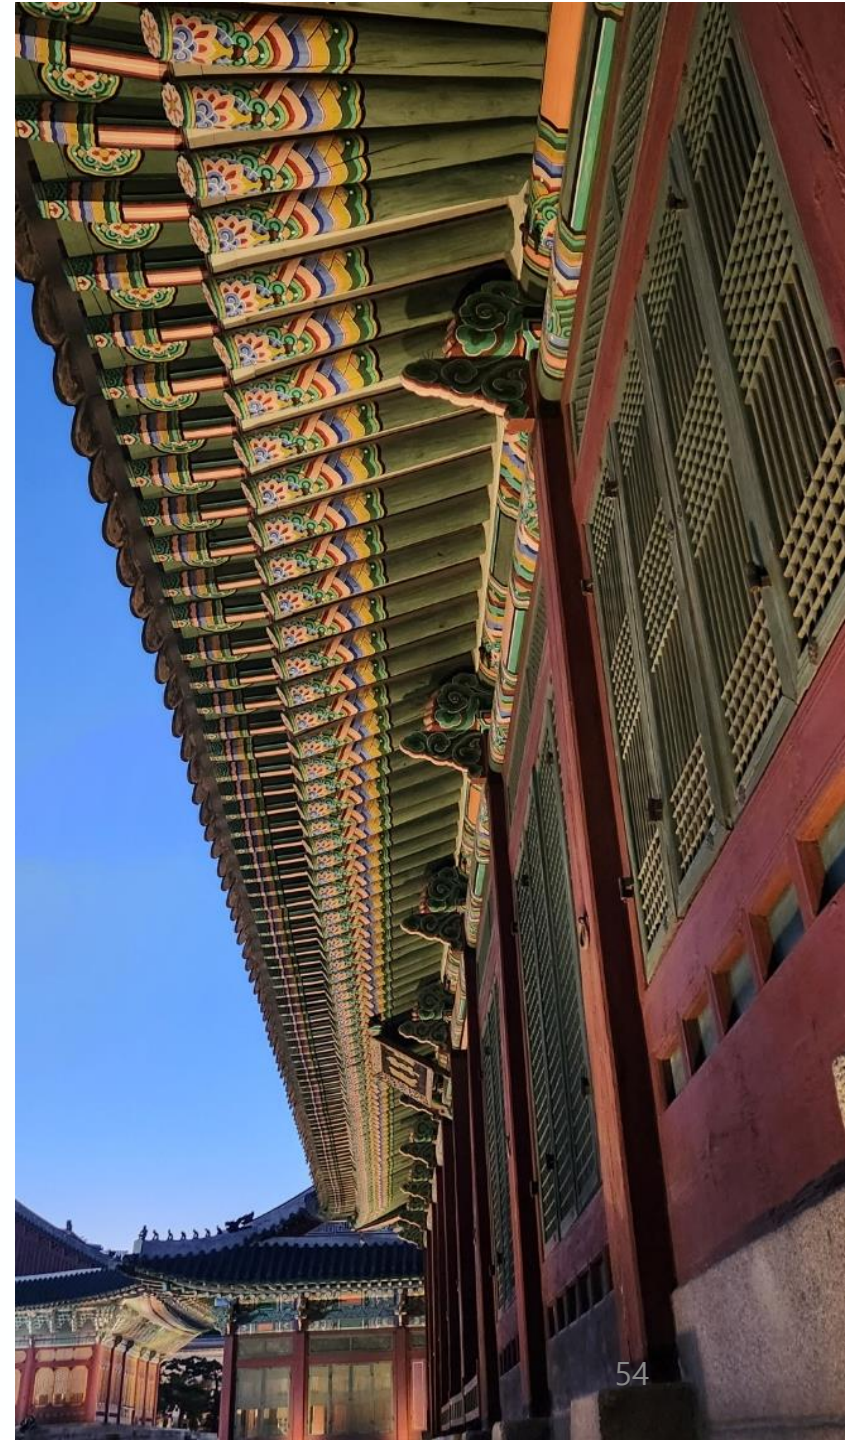


# Calculation for Max File Size with Block Sizes

Blocksize (KByte)	1	2	4	8
Blocks first stage	12	12	12	12
Blocks first indirection	256	512	1,024	2,048
Blocks second indirection	65,536	262,144	1,048,576	4,194,304
Blocks third indirection	16,777,216	134,217,728	1,073,741,824	8,589,934,592
maximum number of entry presented	16,843,020	134,480,396	1,074,791,436	8,594,130,956
maximum file size(Gbytes)	16	257	4,100	65,568

- Assumption: 4 bytes / Entry
- Increasing the block size increases maximum file size quadratically
- With 4KB/block, upto 4.1 TB file size is possible
- With 8KB/block, upto 65TB file size is possible

# Exercise and Problems for Filesystem



# Quiz #1: Removing consecutive repeated characters from string

- This program will read a line from a file and **remove repeated consecutive characters from the string** and print out the updated line. In this program we will not use another string to copy string after removing the characters; we repeat the updating process line by line manner until all the file is read and print out the result.

```
First Run:
Enter any string: AABBCDD
String after removing characaters: ABCD

Second Run:
Enter any string: AAAABBBBCCCCDDDD
String after removing characaters: ABCD

Third Run:
Enter any string: AAA AAA BBB CCCC DDDD
String after removing characaters: A A B C D
```



# Hints to Quiz#1: Read a string and remove repeated characters.

```
#include <stdio.h>

int main()
{
    char str[100];
    int i,j,len,len1;

    /*read string*/
    printf("Enter any string: ");
    gets(str);

    /*calculating length*/
    for(len=0; str[len]!='\0'; len++)
        ;

    /*assign 0 to len1 - length of removed characters*/
    len1=0;
```

```
/*Removing consecutive repeated characters from string*/
for(i=0; i<(len-len1);) {
    if(str[i]==str[i+1]) {
        /*shift all characters*/
        for(j=i;j<(len-len1);j++)
            str[j]=str[j+1];
        len1++;
    } else {
        i++;
    }
}
printf("String after removing characaters: %s\n",str);
return 0;
}
```



# Quiz #2,#3: Read/Write Struct From/to a File

- **Q2: Writing Structure to a File using fwrite**

We can use fwrite() function to easily write a *struct person data* in a file. fwrite() function writes the data to the file stream *in the form of binary data block*.

- **Q3: Reading Structure from a File using fread**

We can use fread() function to easily read a *struct person data* stored in a file. This function reads a block of memory from the given stream.

```
struct person {  
    int id;  
    char fname[20];  
    char lname[20];  
};
```

# Answer to Quiz#2

```
#include <stdio.h>
#include <stdlib.h>

// a struct to be read and written
struct person {
    int id;
    char fname[20];
    char lname[20];
};

int main() {
    FILE* outfile;

    // open file for writing
    outfile = fopen("person.bin", "wb");
    if (outfile == NULL) {
        fprintf(stderr, "\nError opened file\n");
        exit(1);
    }
}
```

```
struct person input1 = { 1, "rohan", "sharma" };

// write struct to file
int flag = 0;
flag = fwrite(&input1, sizeof(struct person), 1, outfile);
if (flag) {
    printf("Contents of the structure written "
           "successfully");
}
else
    printf("Error Writing to File!");

// close file
fclose(outfile);

return 0;
}
```

# Answer to Quiz #3

```
// struct from a file
#include <stdio.h>
#include <stdlib.h>

// struct person with 3 fields
struct person {
    int id;
    char fname[20];
    char lname[20];
};

int main( )
{
    FILE* infile;
    // Open person.dat for reading
    if( (infile = fopen("person1.dat", "r")) == NULL) {
        fprintf(stderr, "WnError opening fileWn");
        exit(1);
    }
```

```
// reading to read_struct
fread(&read_struct, sizeof(read_struct), 1, infile);

printf("Name: %s %s WnID: %d", read_struct.fname read_struct.lname, read_struct.id);

// close file
fclose(infile);

return 0;
}
```

# Quiz #4 Copying a file contents to another file

- Usage of the program in command line  
    <program-name> <source-file-name> <destination-file-name>  
ex) copy dataSrc.1 dataDest.1
  - copy: a executable file name
  - dataSrc.1 : existing file name (binary or text whatever format)
  - dataDest.1: new file name to be generated with the contents of dataSrc.1
- We can use fgetc() function and fputc() function to read and write a byte from a file, respectively.

# Answer to Quiz#4

```
#include <stdio.h>
#include <stdlib.h> // For exit()

int main(int argc ,char **argv)
{
    FILE *fptr1 = stdin, *fptr2 = stdout;
    int c;

    if(argc < 3) {
        printf("Usage: %s <src> <dest>\n", argv[0]);
        exit(1);
    }

    // Open one file for reading
    fptr1 = fopen (argv[1], "r");
    if (fptr1 == NULL) {
        printf("Cannot open file %s\n",argv[1]);
        exit(1);
    }
```

```
    // Open another file for writing
    fptr2 = fopen(argv[2], "w");
    if (fptr2 == NULL) {
        printf("Cannot open file %s\n", argv[2]);
        exit(1);
    }

    // Read contents from file
    while ((c = fgetc(fptr1)) != EOF)
    {
        fputc(c, fptr2);
    }

    printf("Contents copied to %s\n", filename);
    fclose(fptr1);
    fclose(fptr2);
    return 0;
}
```

# Quiz#5 Append Content of one text file to another

- Given the source and destination text files, the task is to append the content from source file to destination file and then display the content of the destination file.
- **Approach:**
  1. Open **file1.txt** and **file2.txt** with "a+"(append and read) option, so that the previous content of the file is not deleted. If files don't exist, they will be created.
  2. Explicitly write a newline ("\n") to the destination file to enhance readability.
  3. Write content from source file to destination file.
  4. Display the contents in file2.txt to console (stdout).

Input:

**file1.txt**

This is line one in file1  
Hello World.

**file2.txt**

This is line one in file2  
Programming is fun.

Output:

This is line one in file2  
Programming is fun.  
This is line one in file1  
Hello World.

# Answer to Quiz#5

```
#include <stdio.h>

// Function that appends the contents
void appendFiles(char source[],
                 char destination[])
{
    // declaring file pointers
    FILE *fp1, *fp2;

    // opening files
    fp1 = fopen(source, "a+");
    fp2 = fopen(destination, "a+");

    // If file is not found then return.
    if (!fp1 && !fp2) {
        printf("Unable to open/"
              "detect file(s)\n");
        return;
    }
```

```
char buf[100];

// explicitly writing "\n"
// to the destination file
// so to enhance readability.
fprintf(fp2, "\n");

// writing the contents of
// source file to destination file.
while (!feof(fp1)) {
    fgets(buf, sizeof(buf), fp1);
    fprintf(fp2, "%s", buf);
}
```

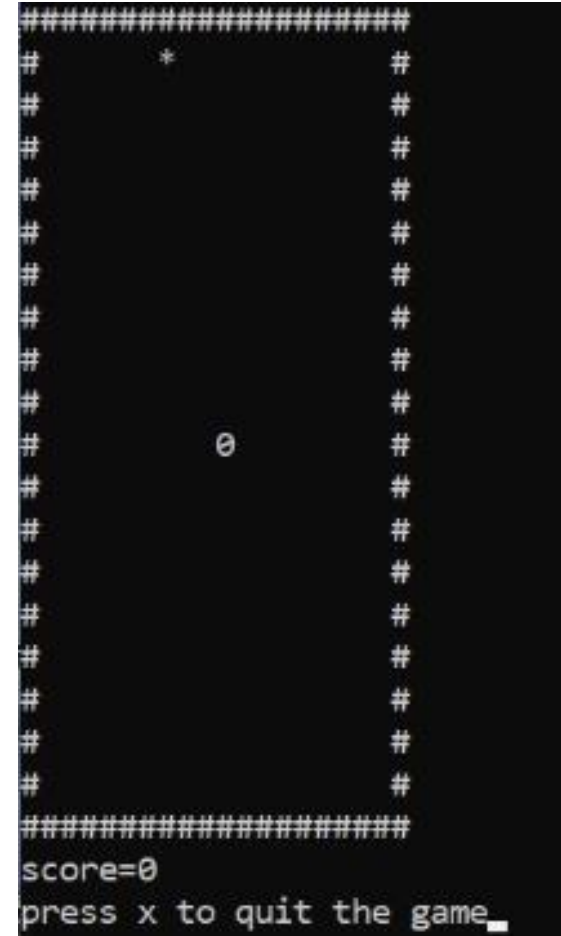
```
// printing the results to stdout
rewind(fp2);
while (!feof(fp2)) {
    fgets(buf, sizeof(buf), fp2);
    printf("%s", buf);
}

int main()
{
    char source[] = "file1.txt",
          destination[] = "file2.tx

    // calling Function with file names.
    appendFiles(source, destination);
    return 0;
}
```

# Problem #1: Snake Game in C

- Below given some functionalities of this game:
  - The snake is represented with a **0**(zero) symbol.
  - The fruit is represented with an **\***(asterisk) symbol.
  - The snake can move in any direction according to the user with the help of the keyboard (**W**, **A**, **S**, **D** keys).
  - When the snake eats a fruit the score will increase by 10 points.
  - The fruit will generate automatically randomly within the boundaries.
  - Whenever the snake will touch the boundary the game is over.





# Ans(1). to Prob-1

```
// C program to build the
// complete
// snake game
#include <conio.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int i, j, height = 20, width = 20;
int gameover, score;
int x, y, fruitx, fruity, flag;
```

```
void main()
{
    int m, n;
    setup();
    while (!gameover) {
        // Function Call
        draw();
        input();
        logic();
    }
}
```

```
// Function to generate the fruit
// within boundary
void setup()
{
    gameover = 0;

    // Stores height and width
    x = height / 2;
    y = width / 2;
label1:
    fruitx = rand() % 20;
    if (fruitx == 0)
        goto label1;
label2:
    fruity = rand() % 20;
    if (fruity == 0)
        goto label2;
    score = 0;
}
```

```
// Function to draw the boundaries
void draw() {
    system("cls");
    for (i = 0; i < height; i++) {
        for (j = 0; j < width; j++) {
            if (i == 0 || i == width - 1 || j == 0
                || j == height - 1) {
                printf("#");
            } else {
                if (i == x && j == y)
                    printf("O");
                else if (i == fruitx && j == fruity)
                    printf("*");
                else
                    printf(" ");
            }
        }
        printf("\n");
    }
    // Print the score after the
    // game ends
    printf("score = %d", score);
    printf("\n");
    printf("press X to quit the game");
}
```

# Ans(2) to Prob-1

```
// Function to take the input
```

```
void input() {  
    if (kbhit()) {  
        switch (getch()) {  
            case 'a':  
                flag = 1;  
                break;  
            case 's':  
                flag = 2;  
                break;  
            case 'd':  
                flag = 3;  
                break;  
            case 'w':  
                flag = 4;  
                break;  
            case 'x':  
                gameover = 1;  
                break;  
        }  
    }  
} // end of input function
```

```
// Function for the logic  
// behind each movement
```

```
void logic()  
{  
    sleep(0.01);  
    switch (flag) {  
        case 1:  
            y--;  
            break;  
        case 2:  
            x++;  
            break;  
        case 3:  
            y++;  
            break;  
        case 4:  
            x--;  
            break;  
        default:  
            break;  
    }  
}
```

```
// If the game is over  
if (x < 0 || x > height  
    || y < 0 || y > width)  
    gameover = 1;
```

```
// If snake reaches the fruit  
// then update the score  
if (x == fruitx && y == fruity) {  
label3:  
    fruitx = rand() % 20;  
    if (fruitx == 0)  
        goto label3;
```


```
// After eating the above fruit  
// generate new fruit  
label4:  
    fruity = rand() % 20;  
    if (fruity == 0)  
        goto label4;  
    score += 10;  
}  
} // end of logic function
```

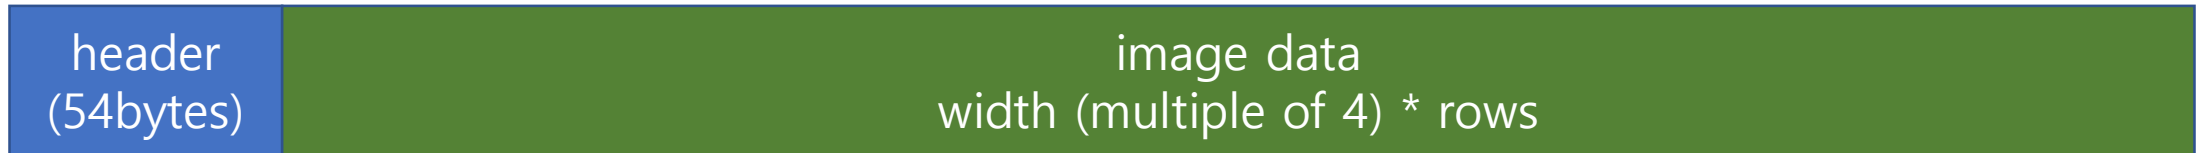
# Problem-2: BMP file

- **Problem**

- 1. Dump BMP file header information: width, rows, and file size**
- 2. Dump the BMP file data as hexadecimal as xxd command do**
- 3. Change the BMP file to make it grayscale**
  - Change Red and Blue value as the same value of Green for all pixels
- 4. Encrypt a message in an BMP image (Steganography)**
  - Use LSB bit of each pixel's blue and red color to store information

# BMP file format

- BMP file format : 54 bytes header + image data
- Each pixel consists of Blue, Green, and Red color (8bits \* 3 = 24bits)
  - each pixel data : 
- Each row data should be multiple of 4 bytes (padding is required for each row data to adjust multiple of 4 byte size per row)



# BMP file format

Component	Size
Header	54 byte
Palette (optional)	0 byte (for 24-bit RGB image)
Image Data	file size – 54 (for 24-bit RGB image)

- header size : 54 bytes fixed
- header begins with a magic#: **0x4d42** (16bits)
- important header fields:
  - size: file size in byte
  - width\_px: number of pixels per row
  - height\_px: number of rows
  - bits\_per\_pixel: number of bits per pixel
    - 24 for BGR colors (8bits \* 3 colors)
  - image\_size\_bytes : the size of image data
    - file size – header size

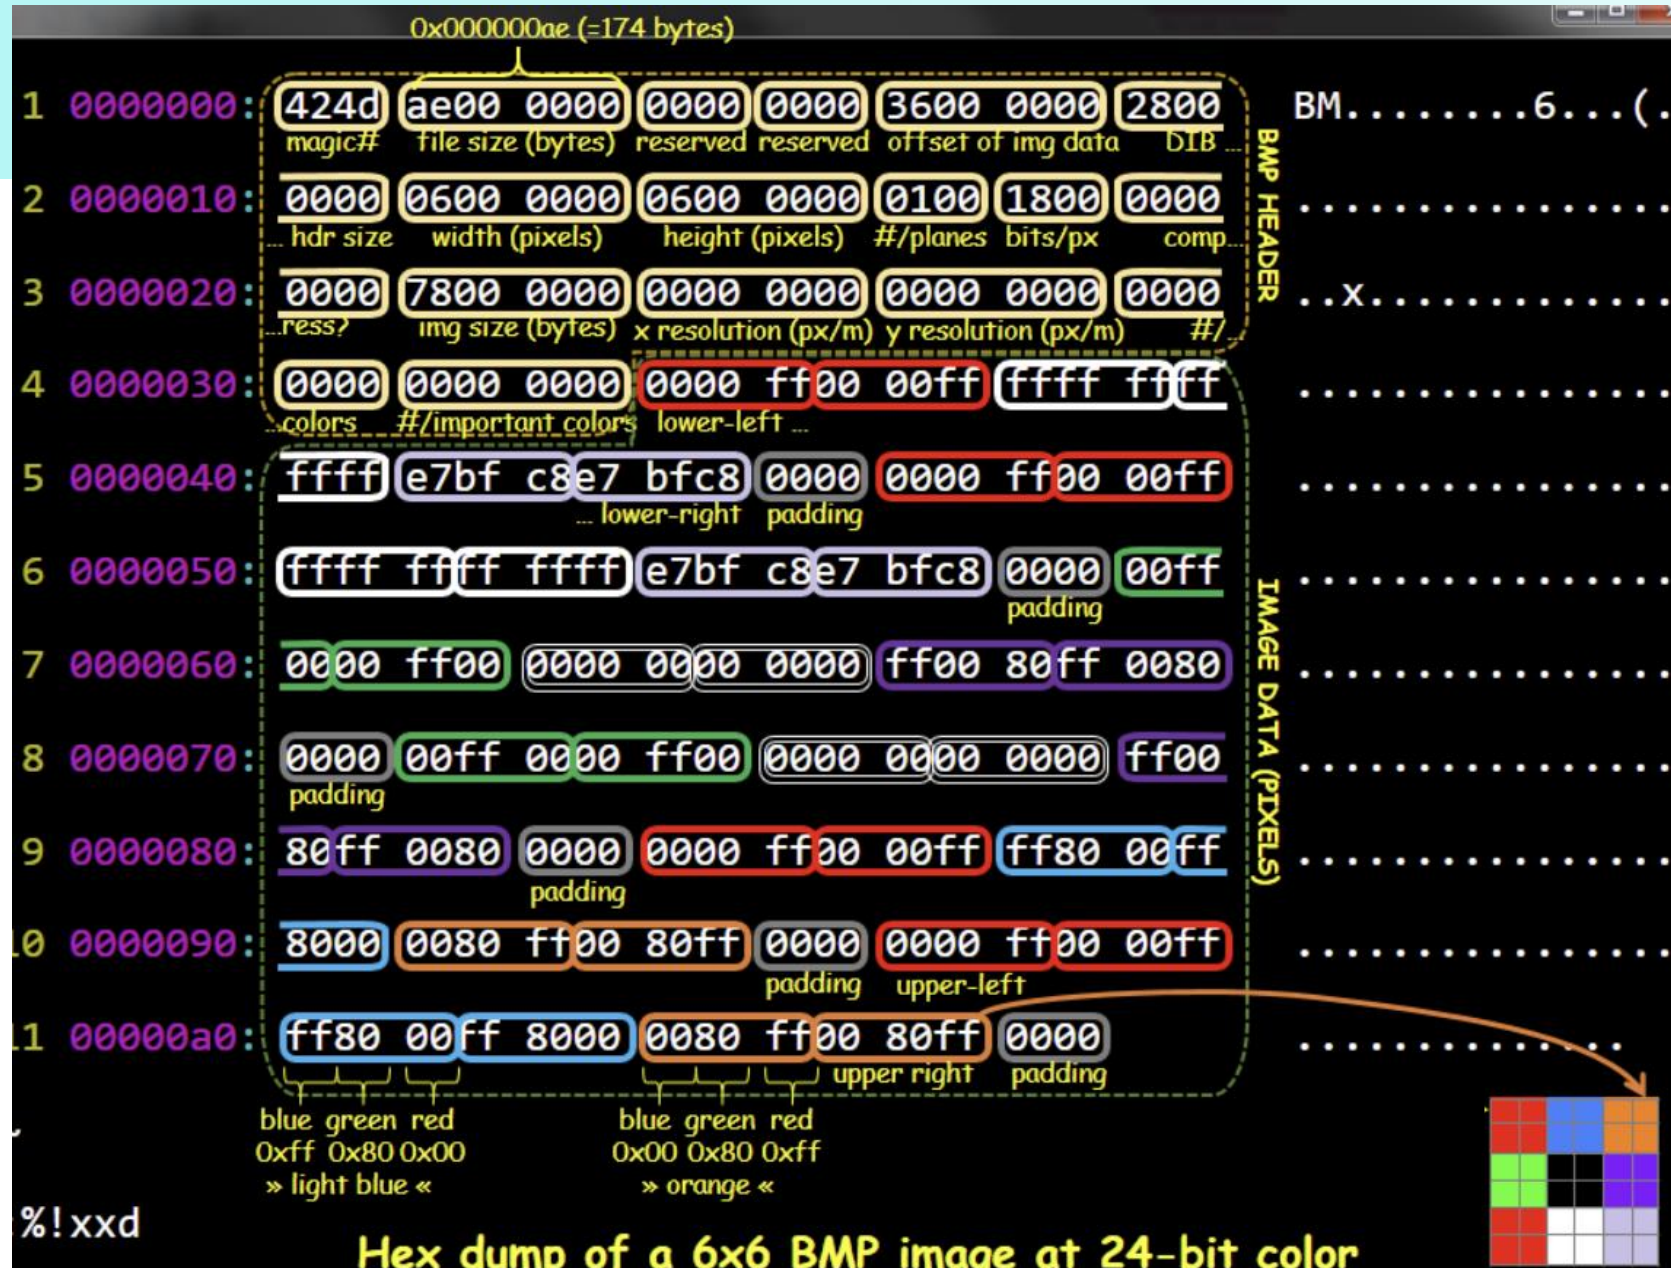
<BMP Header>

```
typedef struct {
    uint16_t    type;
    uint32_t    size;
    uint16_t    reserved1;
    uint16_t    reserved2;
    uint32_t    offset;
    uint32_t    dib_header_size;
    int32_t     width_px;
    int32_t     height_px;
    uint16_t    num_planes;
    uint16_t    bits_per_pixel;
    uint32_t    compression;
    uint32_t    image_size_bytes;
    int32_t     x_resolution_ppm;
    int32_t     y_resolution_ppm;
    uint32_t    num_colors;
    uint32_t    important_colors;
} BMPHeader;
```

# BMP format

\$ xxd <BMP\_file>

- header size = 54 bytes
- image size = 6 pixels x 6 pixels = **36 pixels size**
- Each row = 6 \* 3 bytes = 18 bytes → **20 bytes** (for 4 bytes alignment)
- image data size = 20 bytes/row \* 6 = **120 bytes**
- file size = header size + image data  
= 54 + 120 = **174**





# Endianness of BMP file: Little Endian

- little-endian: LSB comes first (Intel x86, x86-64)
- Big-endian : MSB comes first (IBM 360, Motorola 68000)

```
#include <stdio.h>
// testing endianness with 4 byte value
int main()
{
    FILE *fp = fopen("data","w+");

    unsigned int data = 0x12345678;
    fwrite(&data, sizeof(data), 1, fp);

    fclose(fp);
    return 0;
}
```

yk@peace:~/systemprj/file\$ xxd data  
00000000: 7856 3412

<little-endian order>

03:	78
02:	56
01:	34
00:	12

424d ae00 0000  
magic# file size (bytes)

05:	00
04:	00
03:	00
02:	ae
01:	4d
00:	42

file size → 00 00 00 ae

magic# → 4d 42

- magic#: 0x4d42
- file size: 0x000000ae  
= A\*16+E = 174

BMP file is little-endian format

# bmp.h

#pragma pack(1) directive ensures that the header structure is really 54-byte long by using 1-byte alignment.

#pragma pack(1)

```
typedef struct {
    uint16_t  type;           // Total: 54 bytes
    uint32_t  size;           // Magic identifier: 0x4d42
    uint16_t  reserved1;      // File size in bytes
    uint16_t  reserved2;      // Not used
    uint32_t  offset;         // Not used
    uint32_t  dib_header_size; // Offset to image data in bytes from beginning of file (54)
    int32_t   width_px;       // DIB Header size in bytes (40 bytes)
    int32_t   height_px;      // width of the image
    uint16_t  num_planes;     // Height of image
    uint16_t  bits_per_pixel; // Number of color planes
    uint32_t  compression;    // Bits per pixel
    uint32_t  image_size_bytes; // Compression type
    int32_t   x_resolution_ppm; // Image size in bytes
    int32_t   y_resolution_ppm; // Pixels per meter
    uint32_t  num_colors;      // Pixels per meter
    uint32_t  important_colors; // Number of colors
    uint32_t  important_colors; // Important colors
} BMPHeader;

typedef struct {
    BMPHeader header;
    unsigned char* data;
} BMPImage;
```



# Hint to Prob-2: The Skeleton Code of bmp.c

```
#include <stdio.h>
#include <stdlib.h>
#include "bmp.h"

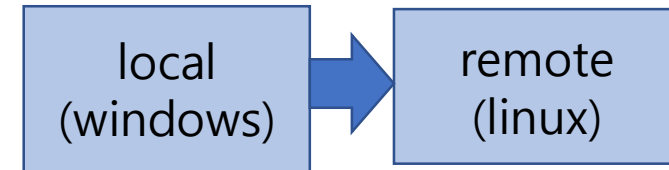
int main (int argc, char *argv[])
{
    BMPImage bmp_img;
    FILE *fp1 = fopen(argv[1], "r");
    FILE *fp2 = fopen(argv[2], "w");

    read_header(fp1, &(bmp_img.header));
    bmp_img.data = read_data (fp1, bmp_img);
    change_color_grayscale(bmp_img.data, bmp_hd.width_px, bmp_hd.height_px, bmp_hd.);
    write_data(fp2, &bmp_img);
    return (1);
}
```

# How to transfer files to/from remote host

- local에서 원격 host로 이미지 파일 보내기(scp)
  - **scp** *<source-file>* *<user\_id>@<remot-host>:<target-file>*

```
C:\Users\yk\Pictures>scp khj.jpg yk@peace.handong.edu:image.jpg
yk@peace.handong.edu's password:
khj.jpg
C:\Users\yk\Pictures>
```



- 원격 host에서 local로 이미지 파일 가져오기(scp)
  - **scp** *<user\_id>@<remot-host>:<source\_file>* *<target-file>*

```
C:\Users\yk\Pictures>scp yk@peace.handong.edu:image.jpg k.jpg
yk@peace.handong.edu's password:
image.jpg
```

